

Lecture 2: OS Interfaces

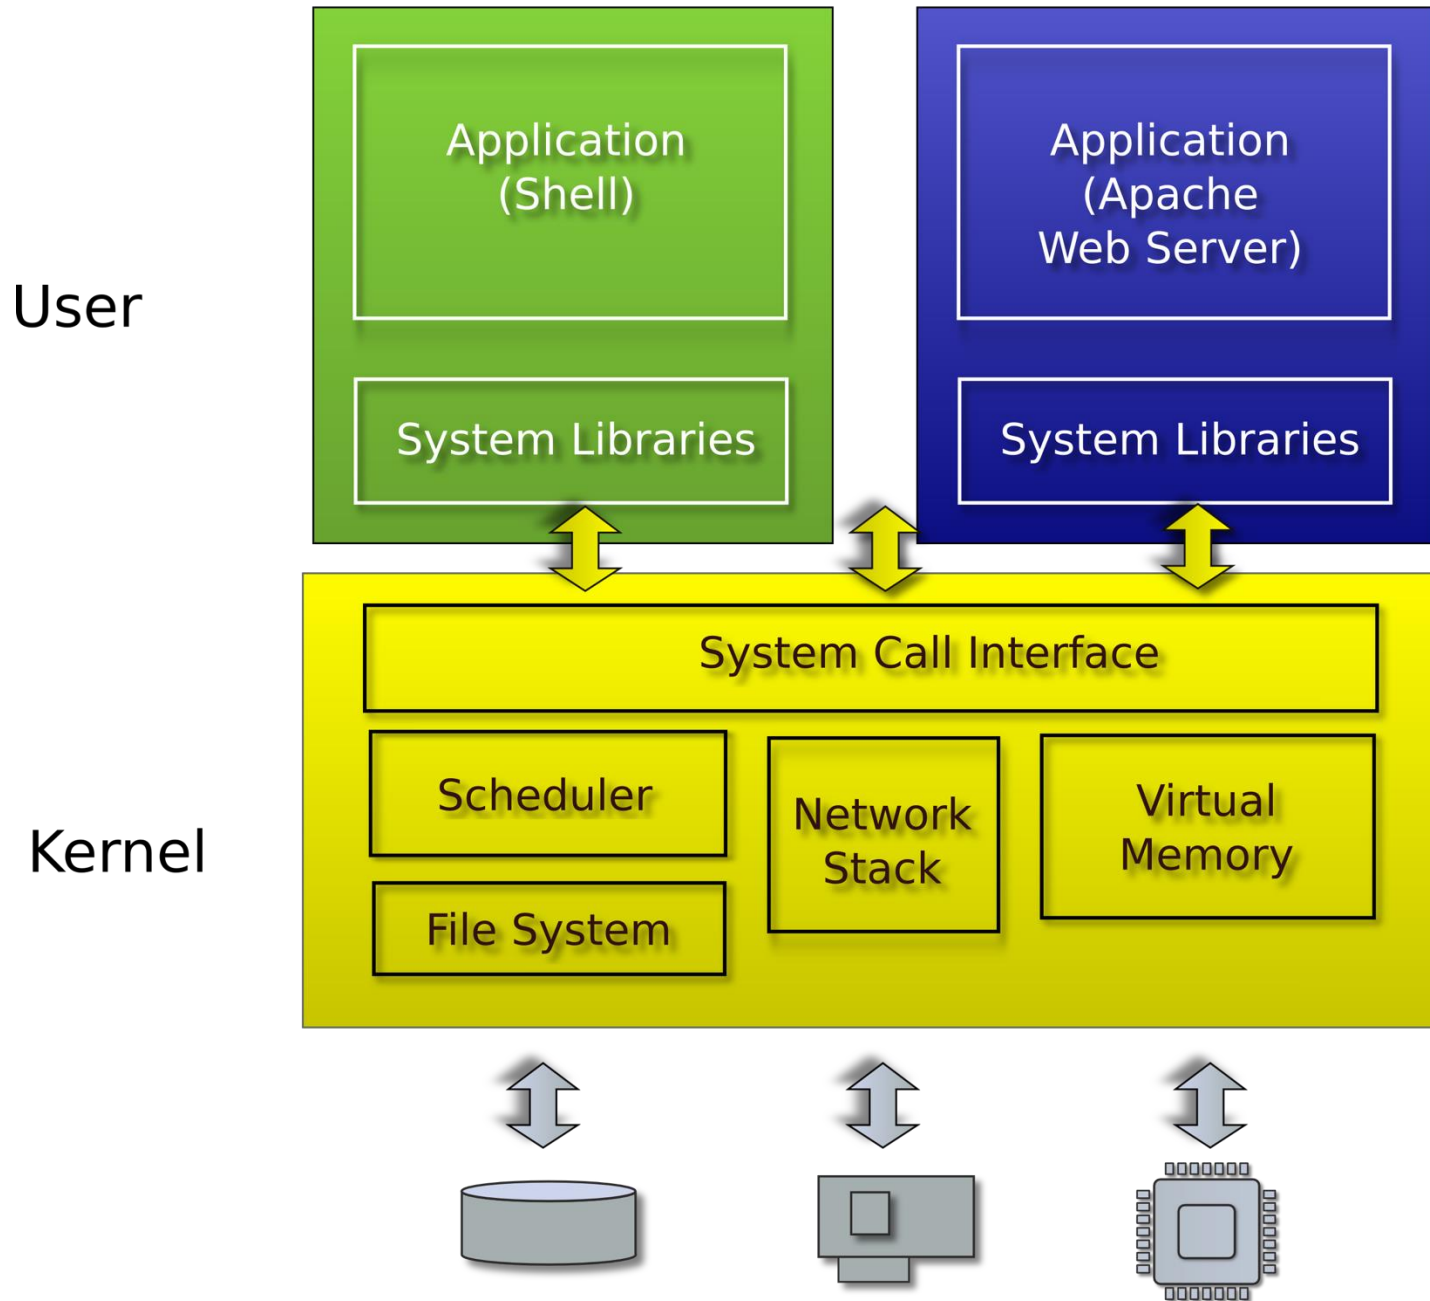
cs5460/6460 Operating Systems

Anton Burtsev
January, 2025

Recap: role of the operating system

- Share hardware across multiple processes
 - Illusion of private CPU, private memory
- Abstract hardware
 - Hide details of specific hardware devices
- Provide services
 - Serve as a library for applications
- Security
 - Isolation of processes
 - Controlled ways to communicate (in a secure manner)

Typical UNIX OS



System calls

- Provide user to kernel communication
 - Effectively an invocation of a kernel function

System calls implement the interface of the OS

Which system calls do we need?

System calls, interface for...

- Processes
 - Creating, exiting, waiting, terminating
- Memory
 - Allocation, deallocation
- Files and folders
 - Opening, reading, writing, closing
- Inter-process communication
 - Pipes

UNIX (and xv6) system calls are designed around **shell**

```
Sun/01.10:/home/aburtsev/projects/xv6-public
```

```
aburtsev-ThinkPad-X1-Carbon-3rd:516-/23:21>ls
```

asm.h	cat.o	entryother.o	fs.o	init.d	kill.d
bio.c	cat.sym	entryother.S	gdbutil	init.o	kill.o
bio.d	console.c	entry.S	_grep*	init.sym	kill.sym
bio.o	console.d	exec.c	grep.asm	ioapic.c	lapic.c
bootasm.d	console.o	exec.d	grep.c	ioapic.d	lapic.d
bootasm.o	cuth*	exec.o	grep.d	ioapic.o	lapic.o
bootasm.S	date.h	fcntl.h	grep.o	kalloc.c	LICENSE
bootblock*	defs.h	file.c	grep.sym	kalloc.d	_ln*
bootblock.asm	dot-bochsrc*	file.d	ide.c	kalloc.o	_ln.asm
bootblock.o*	_echo*	file.h	ide.d	kbd.c	ln.c
bootblockother.o*	echo.asm	file.o	ide.o	kbd.d	ln.d
bootmain.c	echo.c	_forktest*	_init*	kbd.h	ln.o
bootmain.d	echo.d	forktest.asm	init.asm	kbd.o	ln.sym
bootmain.o	echo.o	forktest.c	init.c	kernel*	log.c
buf.h	echo.sym	forktest.d	initcode*	kernel.asm	log.d
BUGS	elf.h	forktest.o	initcode.asm	kernel.ld	log.o
_cat*	entry.o	fs.c	initcode.d	kernel.sym	ls*
cat.asm	entryother*	fs.d	initcode.o	_kill*	_ls.asm
cat.c	entryother.asm	fs.h	initcode.out*	kill.asm	ls.c
cat.d	entryother.d	fs.img	initcode.S	kill.c	ls.d


```
Sun/01.10:/home/aburtsev/projects/xv6-public
```

```
aburtsev-ThinkPad-X1-Carbon-3rd:517-/23:22>
```

Why shell?



[Ken Thompson](#) (sitting) and [Dennis Ritchie](#) (standing) are working together on a [PDP-11](#) (around 1970). They are using Teletype Model 33 terminals.



DEC LA36 DECwriter II Terminal



DEC VT100 terminal, 1980

Suddenly this makes sense

- List all files

```
\> ls
```

```
total 9212
```

```
drwxrwxr-x 3 aburtsev aburtsev 12288 Oct 1 08:27 ./
drwxrwxr-x 43 aburtsev aburtsev 4096 Oct 1 08:25 ../
-rw-rw-r-- 1 aburtsev aburtsev 936 Oct 1 08:26 asm.h
-rw-rw-r-- 1 aburtsev aburtsev 3397 Oct 1 08:26 bio.c
-rw-rw-r-- 1 aburtsev aburtsev 100 Oct 1 08:26 bio.d
-rw-rw-r-- 1 aburtsev aburtsev 6416 Oct 1 08:26 bio.o
```

```
...
```

- Count number of lines in a file (ls.c implements ls)

```
\> wc -l ls.c
```

```
85 ls.c
```

But what is **shell**?

But what is **shell**?

- Normal process
 - Kernel starts it for each user that logs into the system
 - In xv6 shell is created after the kernel boots
- Shell interacts with the kernel through system calls
 - E.g., starts other processes

What happens underneath?

```
\> wc -l ls.c
```

```
85 ls.c
```

```
\>
```

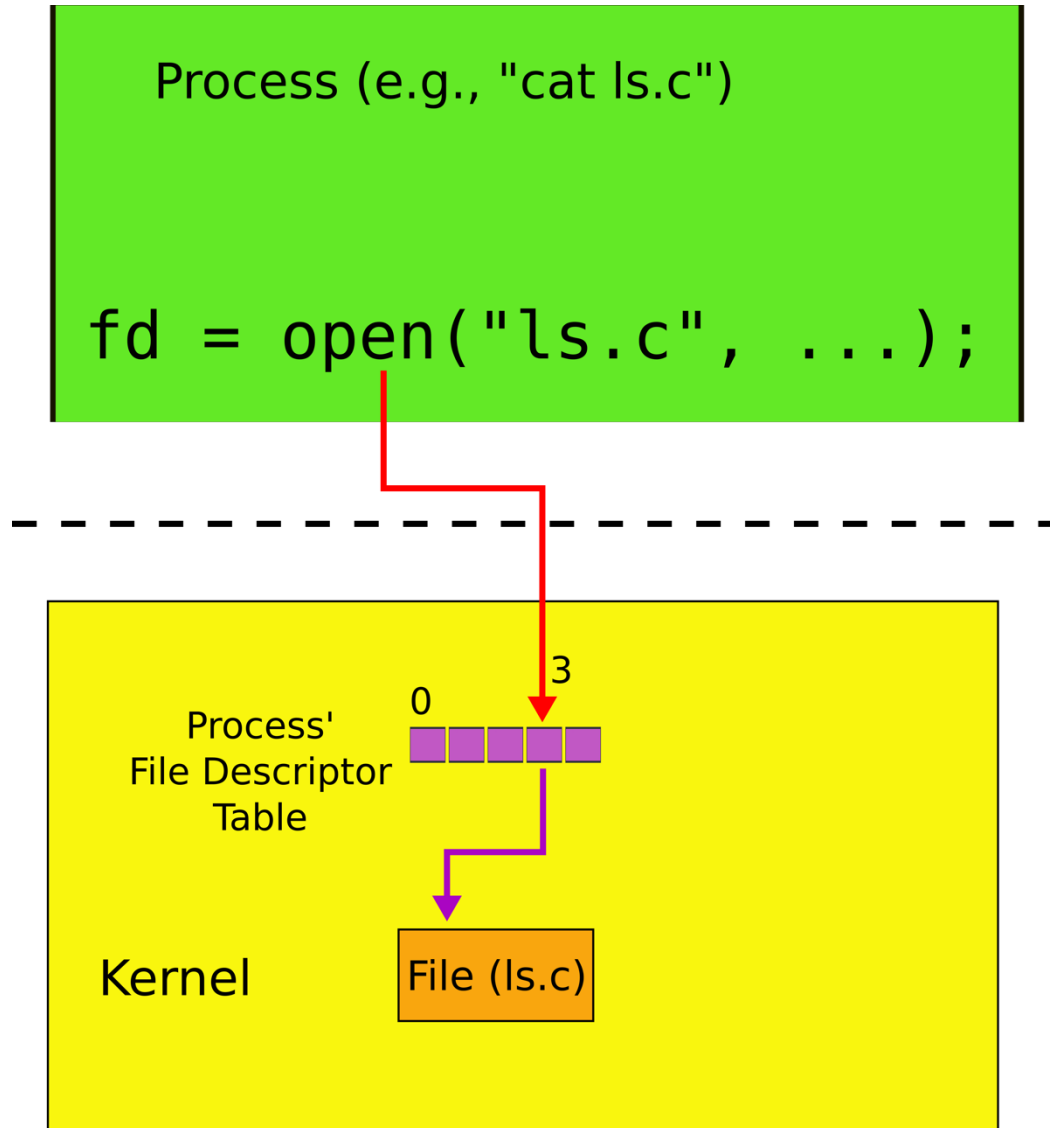
- Shell starts `wc`
 - Creates a new process to run `wc`
 - Passes the arguments (`-l` and `ls.c`)
- `wc` sends its output to the terminal (console)
 - Exits when done with `exit()`
- Shell detects that `wc` is done (`wait()`)
 - Prints (to the same terminal) its command prompt
 - Ready to execute the next command

Console and file I/O

File open

- `fd = open("ls.c", O_RDONLY)`
 - Open a file
 - Operating system returns a file descriptor

File descriptors



File descriptors

- An index into a table, i.e., just an integer
- The table maintains pointers to “file” objects
 - Abstracts files, devices, pipes
 - In UNIX **everything is a file** – all objects provide file interface
- Process may obtain file descriptors through
 - Opening a file, directory, device
 - By creating a pipe
 - Duplicating an existing descriptor

File I/O

- `fd = open("foobar.txt", O_RDONLY)` – open a file
 - Operating system returns a file descriptor
- `read(fd, buf, n)` – read `n` bytes from `fd` into `buf`
- `write(fd, buf, n)` – write `n` bytes from `buf` into `fd`

File descriptors: two processes


Process (e.g., "cat ls.c")

```
read(3, buf, size);
```

Process (e.g., "wc -l wc.c")

```
read(4, buf, size);
```


Green Process'
File Descriptor
Table



Kernel

File (ls.c)

Blue Process'
File Descriptor
Table



Kernel

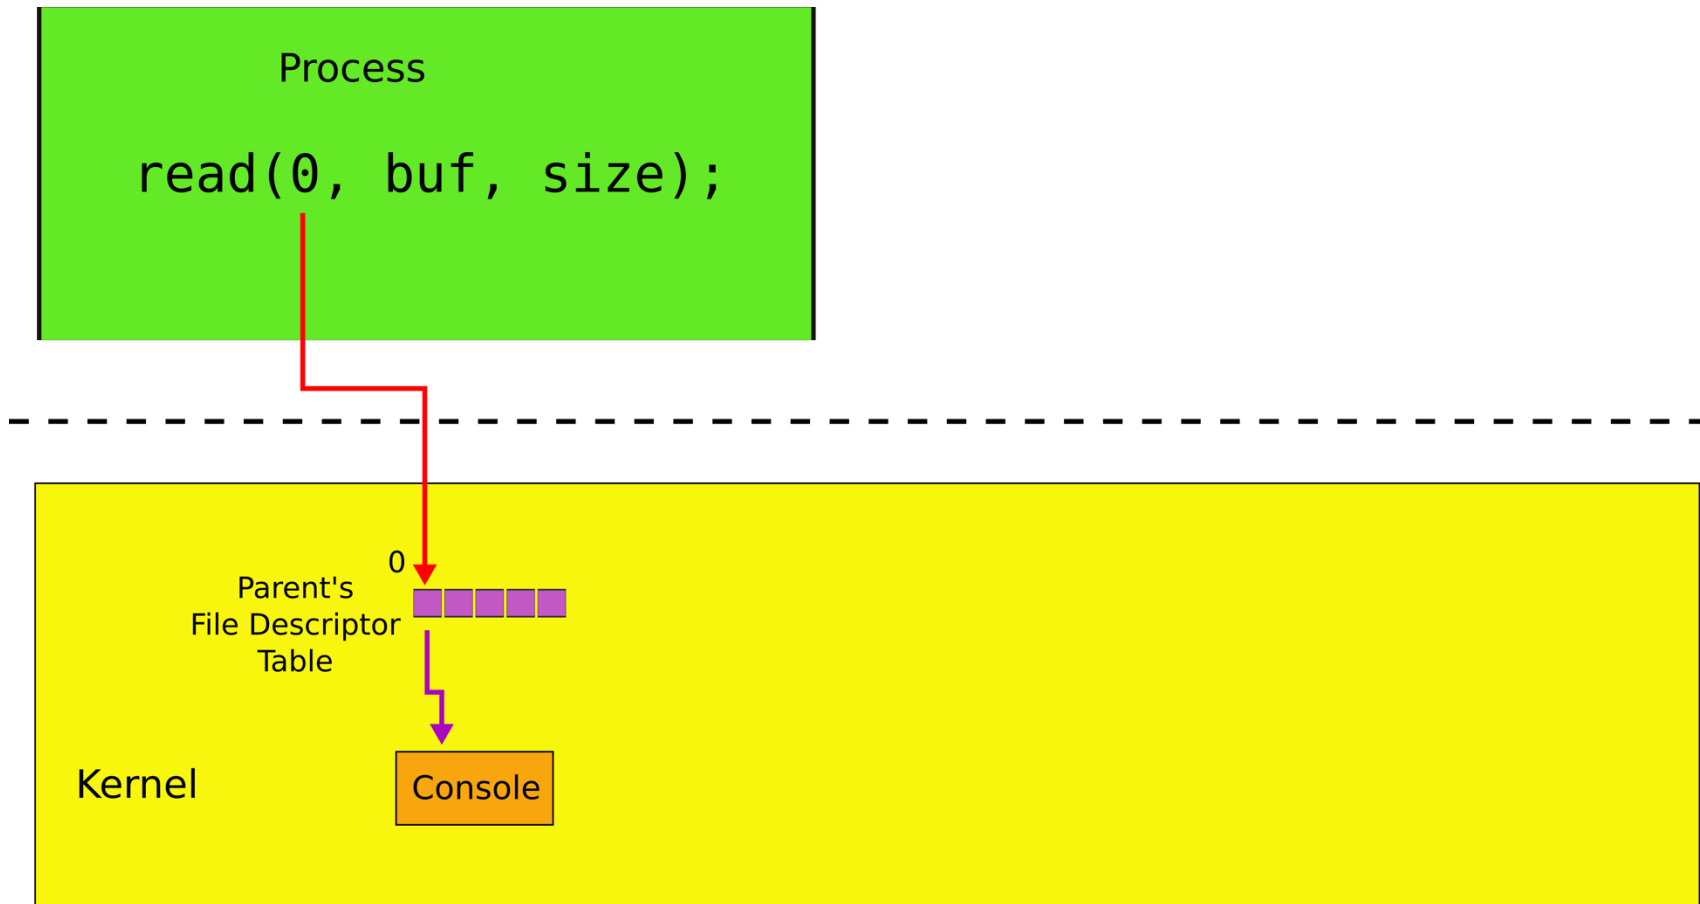
File (wc.c)

Console I/O

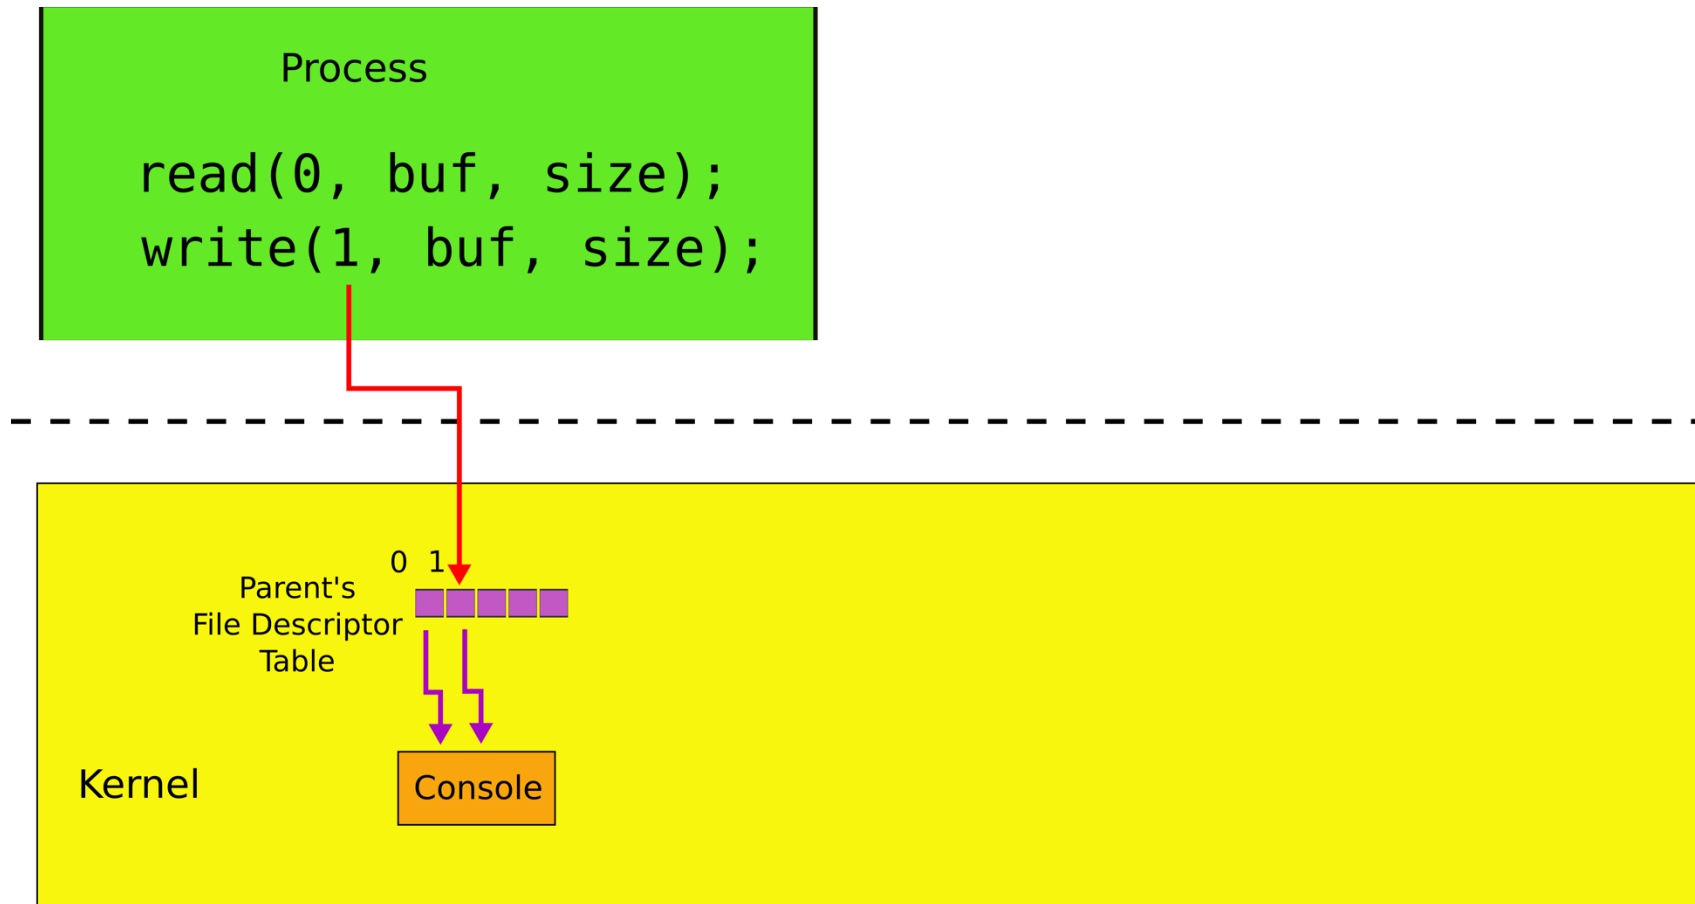
Each process has standard file descriptors

- Numbers are just a convention
 - 0 – standard input
 - 1 – standard output
 - 2 – standard error
- This convention is used by the shell to implement I/O redirection and pipes

Console read (read of standard input)



Console write (write of standard output)



Example: cat

```
1.  char buf[512];
2.  int n;
3.  for(;;) {
4.      n = read(0, buf, sizeof buf);
5.      if(n == 0)
6.          break;
7.      if(n < 0) {
8.          fprintf(2, "read error\n");
9.          exit(); }
10.     if(write(1, buf, n) != n) {
11.         fprintf(2, "write error\n");
12.         exit();
13.     }
14. }
```

Creating processes

fork()

Shell

```
pid = fork()
```

Kernel

fork()

Shell (parent)

32 = fork()

Shell (child)

0 = fork()

Kernel

fork() -- creates a new process

```
1. int pid;  
2. pid = fork();  
3. if(pid > 0){  
4.     printf("parent: child=%d\n", pid);  
5.     pid = wait();  
6.     printf("child %d is done\n", pid);  
7. } else if(pid == 0){  
8.     printf("child: exiting\n");  
9.     exit();  
10. } else {  
11.     printf("fork error\n");  
12. }
```

This is weird... `fork()` creates copies
of the same process, why?

fork() is used together with exec()

- `exec()` -- replaces memory of a current process with a memory image (of a program) loaded from a file

```
char *argv[3];  
argv[0] = "echo";  
argv[1] = "hello";  
argv[2] = 0;  
exec("/bin/echo", argv);  
printf("exec error\n");
```


fork() and exec()

Parent (Shell)

```
32 = fork()
```

Child (Shell)

```
0 = fork();  
exec("/bin/wc", argv);
```

Kernel

fork() and exec()

Parent (Shell)

```
32 = fork()
```

```
main() {  
    ...  
}
```

wc

Kernel

- Still weird... why first `fork()` and then `exec()`?
- Why not `exec()` directly?

I/O Redirection

Motivating example #1

- Normally `wc` sends its output to the console (screen)
 - Count the number of lines in `ls.c`

```
\> wc -l ls.c
```

```
85 ls.c
```

- What if we want to save the number of lines into a file?

Motivating example #1

- Normally `wc` sends its output to the console (screen)
 - Count the number of lines in `ls.c`

```
\> wc -l ls.c
```

```
85 ls.c
```

- What if we want to save the number of lines into a file?
 - We can add an argument

```
\> wc -l ls.c -o foobar.txt
```

Motivating example #1

```
\> wc -l ls.c -o foobar.txt
```

- But there is a better way

```
\> wc -l ls.c > foobar.txt
```

I/O redirection

- **>** redirect output
 - Redirect output of a command into a file

```
\> wc -l ls.c > foobar.txt
```

```
\> cat ls.c > ls-new.c
```

- **<** redirect input
 - Redirect input to read from a file

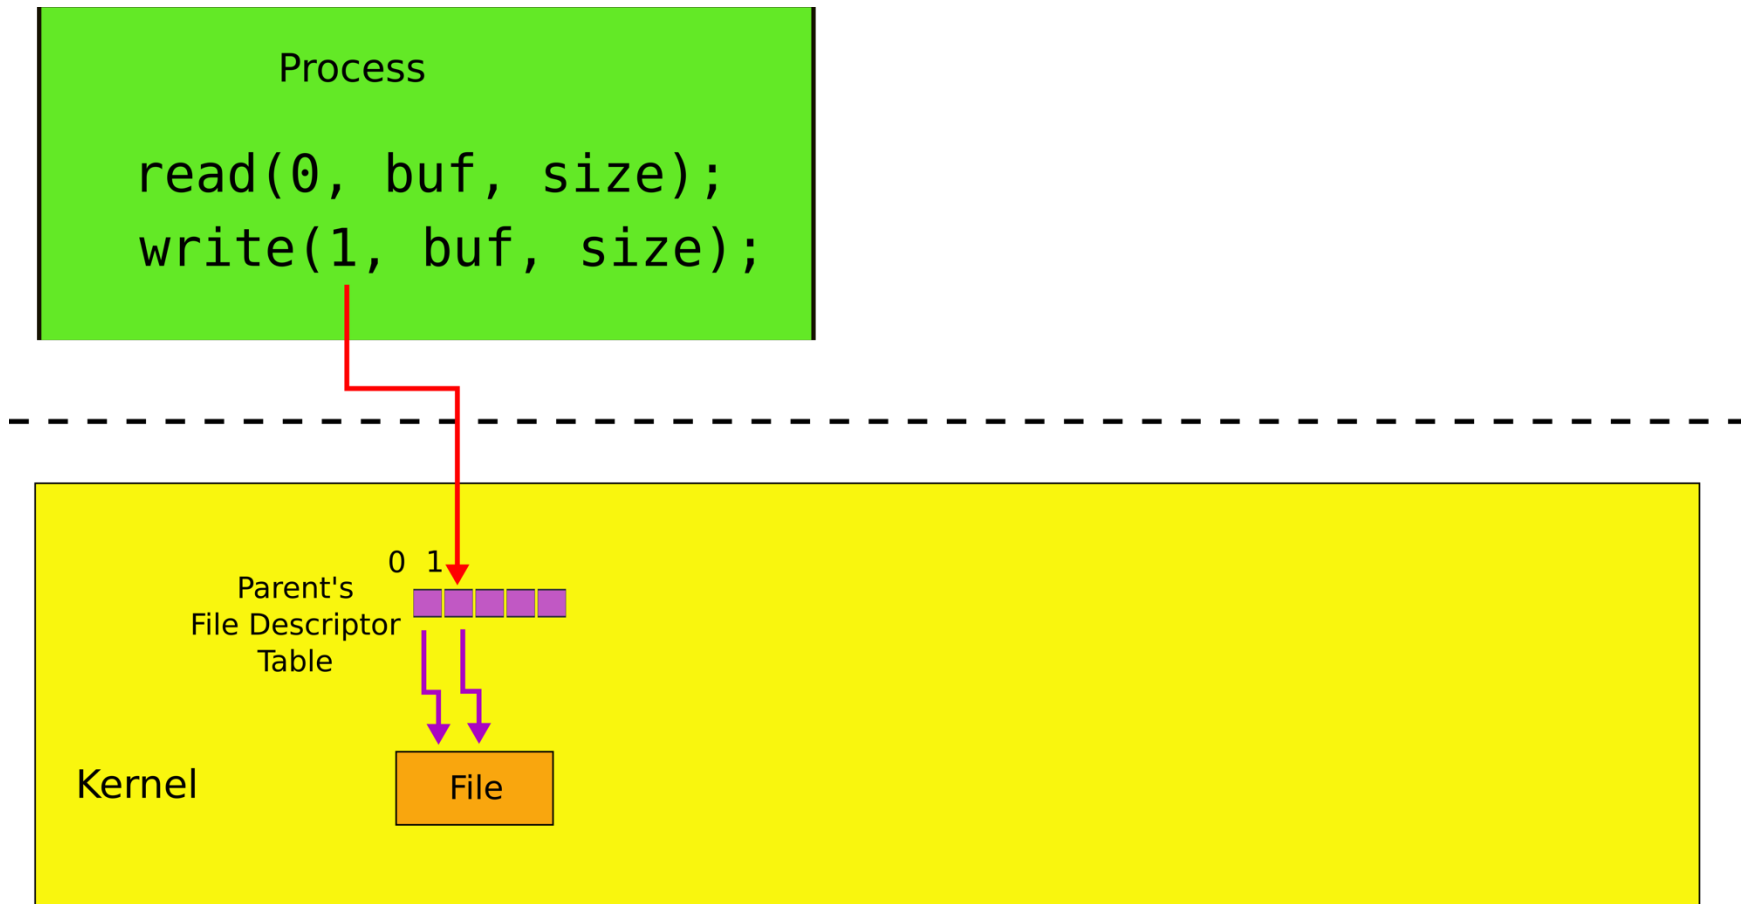
```
\> wc -l < ls.c
```

```
\> cat < ls.c
```

- You can redirect both

```
\> wc -l < ls.c > foobar.txt
```


Standard output is now a file



Powerful design choice

- File descriptors don't have to point to files *only*
 - Any object with the same read/write interface is ok
 - Files
 - Devices
 - Console
 - Pipes

Example: cat

```
1. char buf[512]; int n;  
2. for(;;) {  
3.     n = read(0, buf, sizeof buf);  
4.     if(n == 0)  
5.         break;  
6.     if(n < 0) {  
7.         fprintf(2, "read error\n");  
         exit(); }  
1.     if(write(1, buf, n) != n) {  
2.         fprintf(2, "write error\n");  
3.         exit();  
4.     }  
5. }
```

Why do we need I/O redirection?

Motivating example #2

- We want to see how many strings in `ls.c` contain “`main`”

Motivating example #2

- We want to see how many strings in ls.c contain “main”
 - Imagine we have `grep`
 - `grep` filters strings matching a pattern

```
\>grep "main" ls.c
```

```
main(int argc, char *argv[])
```

- Or the same written differently

```
\>grep "main" < ls.c
```

```
main(int argc, char *argv[])
```

Motivating example #2

- Now we have
 - `grep`
 - Filters strings matching a pattern
 - `wc -l`
 - Counts lines
- Can we combine them?

Pipes

- Imagine we have a way to redirect output of one process into input of another

```
\> cat ls.c | grep main
```

- | (a “pipe”) does redirection

Pipes

- In our example:

```
\> cat ls.c | grep main
```

- `cat` outputs `ls.c` to its output
 - `cat's` output is connected to `grep's` input with the pipe
 - `grep` filters lines that match a specific criteria, i.e., once that have “main”

pipe - inter-process communication

- Pipe is a kernel buffer exposed as a pair of file descriptors
 - One for reading, one for writing
- Pipes allow processes to communicate
 - Send messages to each other

Two file descriptors pointing to a pipe


Process (e.g., "cat ls.c")

```
write(3, buf, size);
```

Process (e.g., "grep main")

```
read(4, buf, size);
```

Green Process'
File Descriptor
Table



Kernel



Pipe

Pipes allow us to connect programs,
i.e., the output of one program to the input of
another

Composability

- Now if we want to see how many strings in ls.c contain “main” we do:

```
\> cat ls.c | grep main | wc -l
```

1

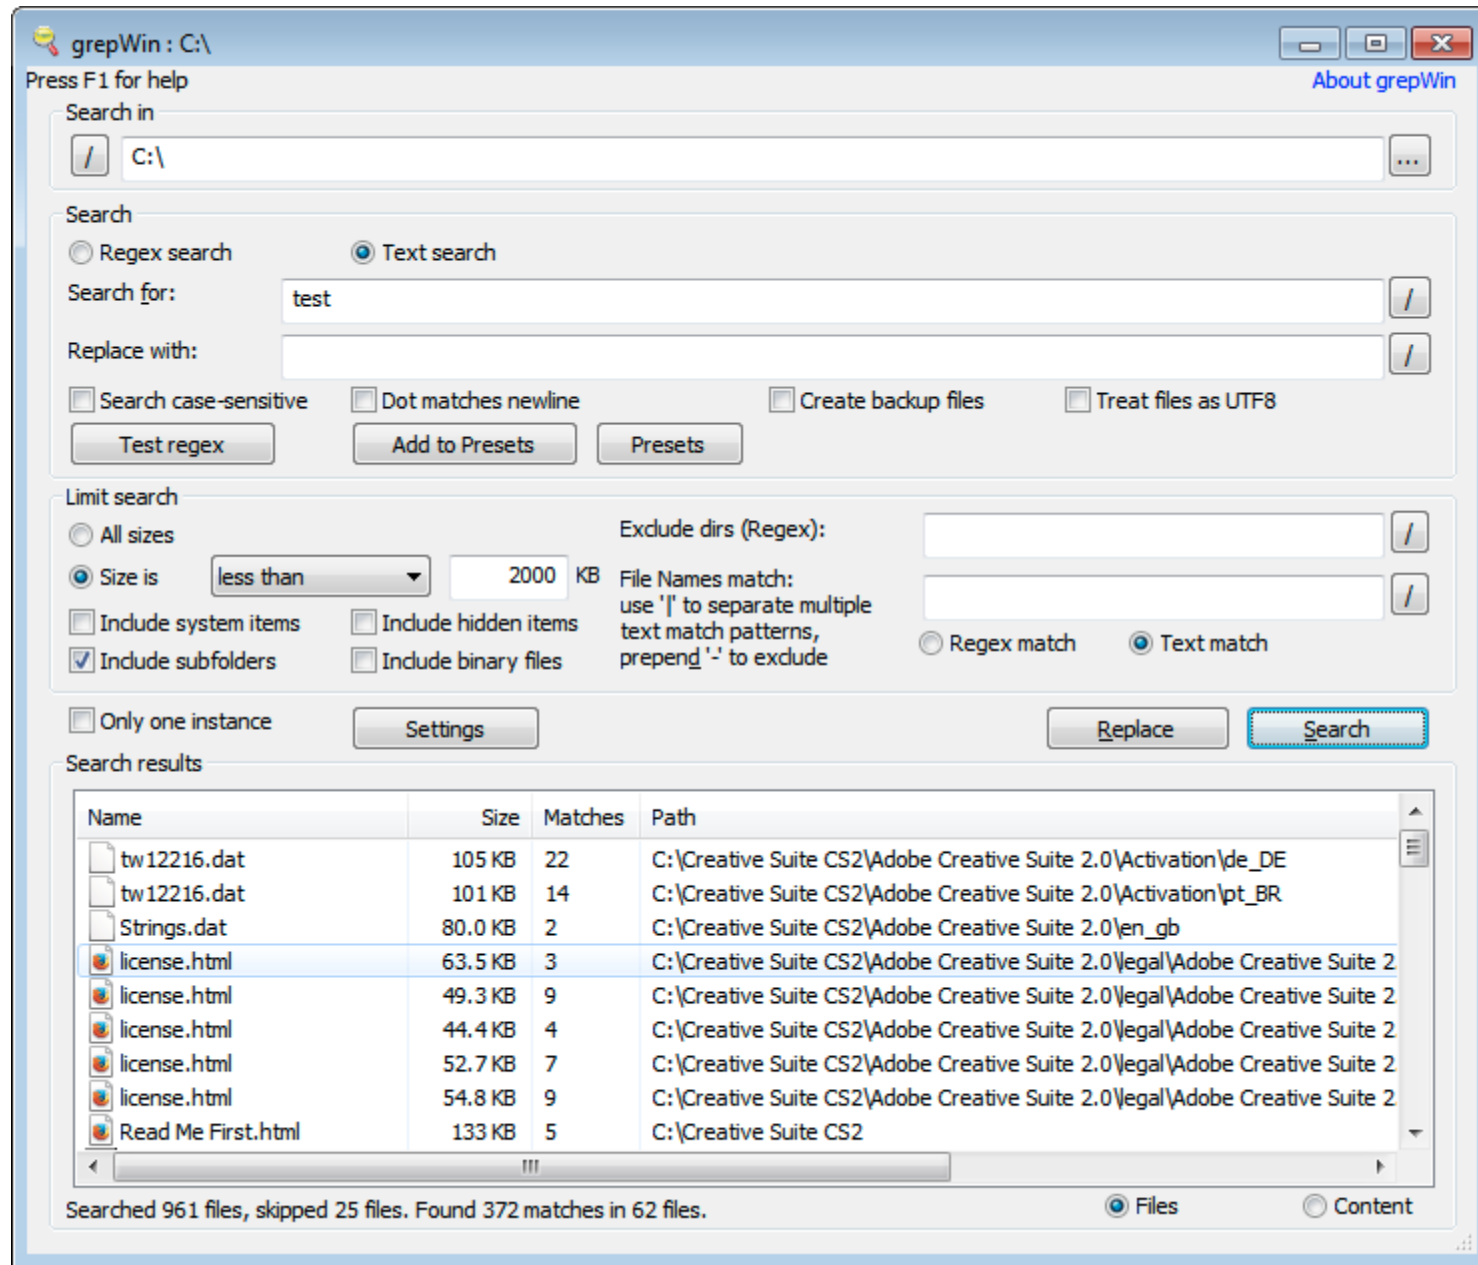
- .. but if we want to count the once that contain “a”:

```
cat ls.c | grep a | wc -l
```

33

- We change only input to grep!
 - Small set of tools (`ls`, `grep`, `wc`) compose into complex workflows

Better than this...



Building I/O redirection

How can we build this?

```
\> cat ls.c | grep main | wc -l
```

- `wc` has to operate on the output of `grep`
- `grep` operates on the output of `cat`

Back to fork()

Shell

```
pid = fork()
```

Kernel

fork()

Shell (parent)

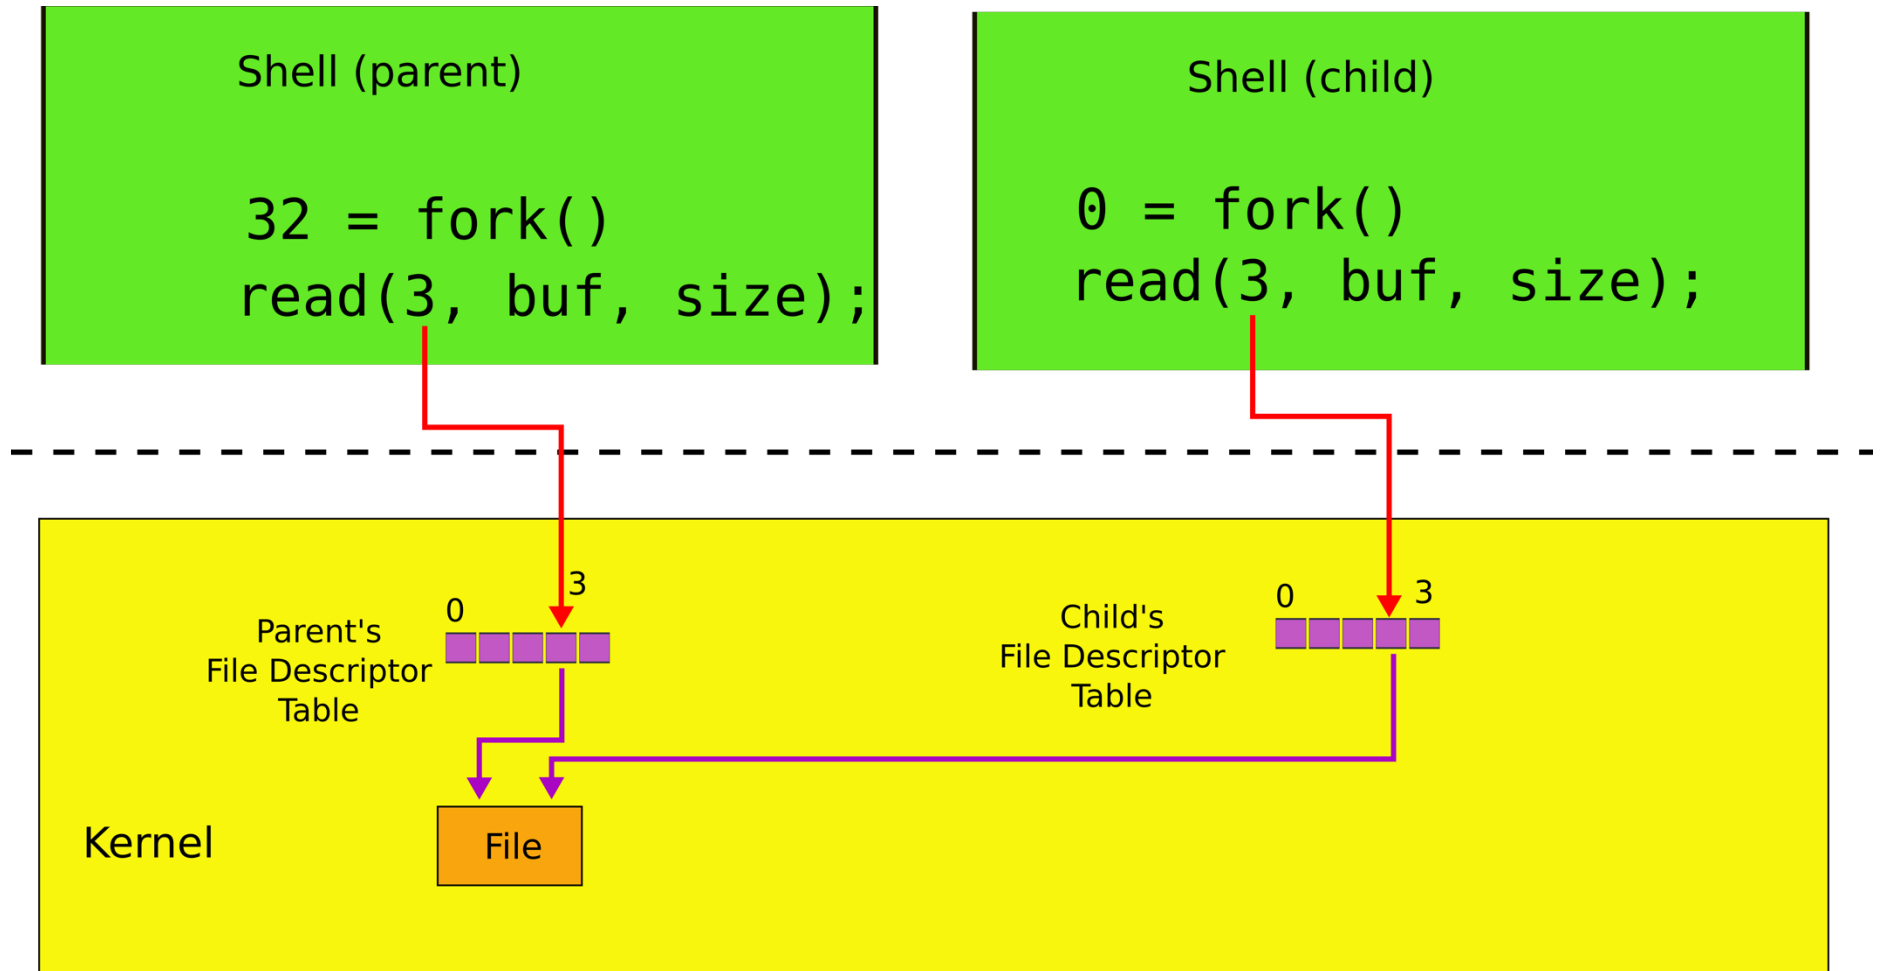
32 = fork()

Shell (child)

0 = fork()

Kernel

File descriptors after fork()




Two system calls for I/O redirection

- `close(fd)` – closes file descriptor
 - **The next opened file descriptor will have the lowest number**

File descriptors after close()/open()

Example: `\> cat < ls.c`

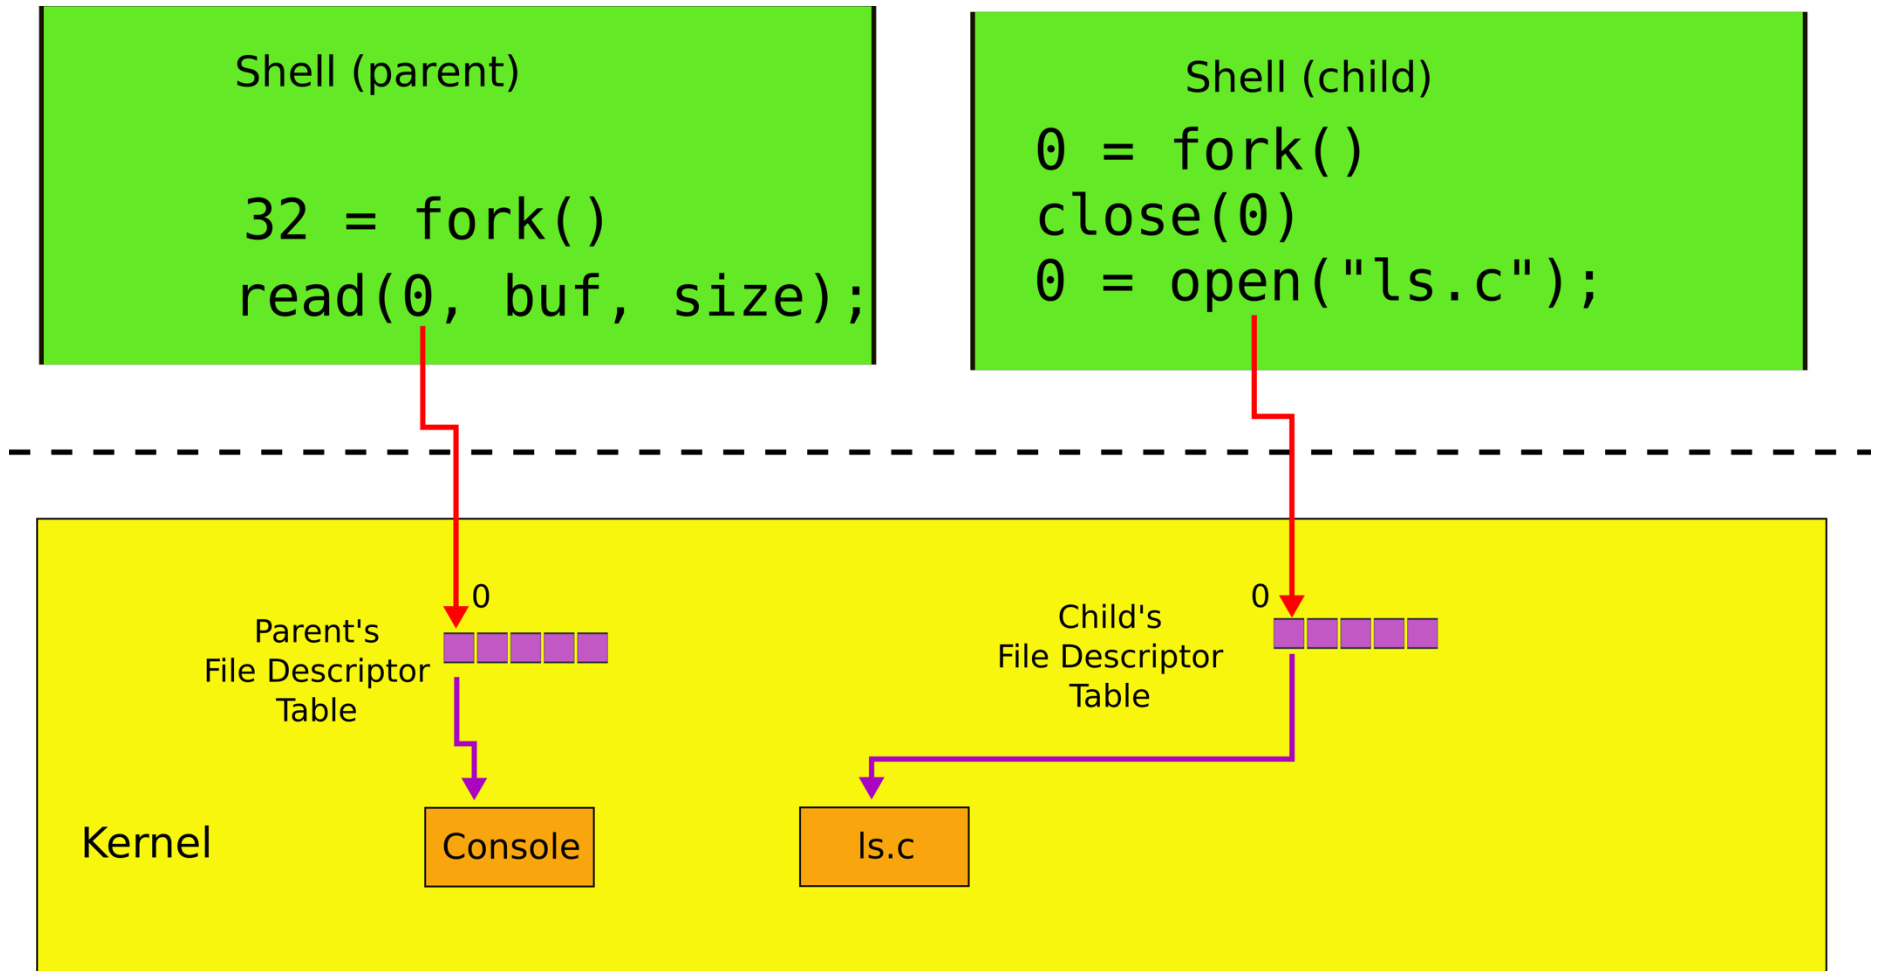


Two system calls for I/O redirection

- `close(fd)` – closes file descriptor
 - The next opened file descriptor will have the lowest number
- `exec()` replaces process memory, but
 - leaves its file descriptor table **intact**
 - A process can create a copy of itself with `fork()`
 - Change the file descriptors for the next program it is about to run
 - And then execute the program with `exec()`

File descriptors after exec()

Example: `\> cat < ls.c`



Example: `\> cat < ls.c`

```
1.     char *argv[2];
2.     argv[0] = "cat";
3.     argv[1] = 0;
4.     if(fork() == 0) {
5.         close(0);
6.         open("ls.c", O_RDONLY);
7.         exec("cat", argv);
8.     }
9.     ...
```

- Poll time

- Inside the `cat` process which file descriptor 0 points to?
- Do we reach line 9?

Why `fork()` not just `exec()`

- The reason for the pair of `fork()/exec()`
 - Shell can manipulate the new process (the copy created by `fork()`)
 - Before running it with `exec()`

Back to Motivating example #2

```
(\> cat ls.c | grep main | wc -l)
```

Pipes

- We now understand how to use a pipe to connect two programs
 - Create a pipe
 - Fork
 - Attach one end to standard output
 - of the left side of “|”
 - Another to the standard input
 - of the right side of “|”

```
1. int p[2];
2. char *argv[2];
3. argv[0] = "wc";
4. argv[1] = 0;
5. pipe(p);
6. if(fork() == 0) {
7.     close(0);
8.     dup(p[0]);
9.     close(p[0]);
10.    close(p[1]);
11.    exec("/bin/wc", argv);
12. } else {
13.    write(p[1], "hello world\n", 12);
14.    close(p[0]);
15.    close(p[1]);
16. }
```

wc on the
read end of
the pipe

Parent

```
write(p[1],  
"hello world\n", 12);
```

wc -l

```
exec("/bin/wc", argv)  
read(0, buf, size);
```

Parent's
File Descriptor
Table



Child's
File Descriptor
Table



Kernel

Pipe

Powerful conclusion

- `fork()`, standard file descriptors, pipes and `exec()` allow complex programs out of simple tools
- They form the core of the UNIX interface

More system calls

Process management

- `exit()` -- terminate current process
- `wait()` -- wait for the child to exit
 - Any child (can be multiple)
 - Return it's process id (`pid`)

Creating files

- `mkdir()` – creates a directory
- `open(..., O_CREATE)` – creates a file
- `mknod()` – creates an empty file marked as device
 - Major and minor numbers uniquely identify the device in the kernel
- `fstat()` – retrieve information about a file

Links, inodes

- Same file can have multiple names – links
 - But unique **inode** number
- **link()** – create a link
- **unlink()** – delete file
- Example, create a temporary file

```
fd = open("/tmp/xyz", O_CREATE|O_RDWR);  
unlink("/tmp/xyz");
```

Xv6 system calls

fork() Create a process

exit() Terminate the current process

wait() Wait for a child process to exit

kill(pid) Terminate process pid

getpid() Return the current process's pid

sleep(n) Sleep for n clock ticks

exec(filename, *argv) Load a file and execute it

sbrk(n) Grow process's memory by n bytes

open(filename, flags) Open a file; the flags indicate read/write

read(fd, buf, n) Read n bytes from an open file into buf

write(fd, buf, n) Write n bytes to an open file

close(fd) Release open file fd

dup(fd) Duplicate fd

pipe(p) Create a pipe and return fd's in p

chdir(dirname) Change the current directory

mkdir(dirname) Create a new directory

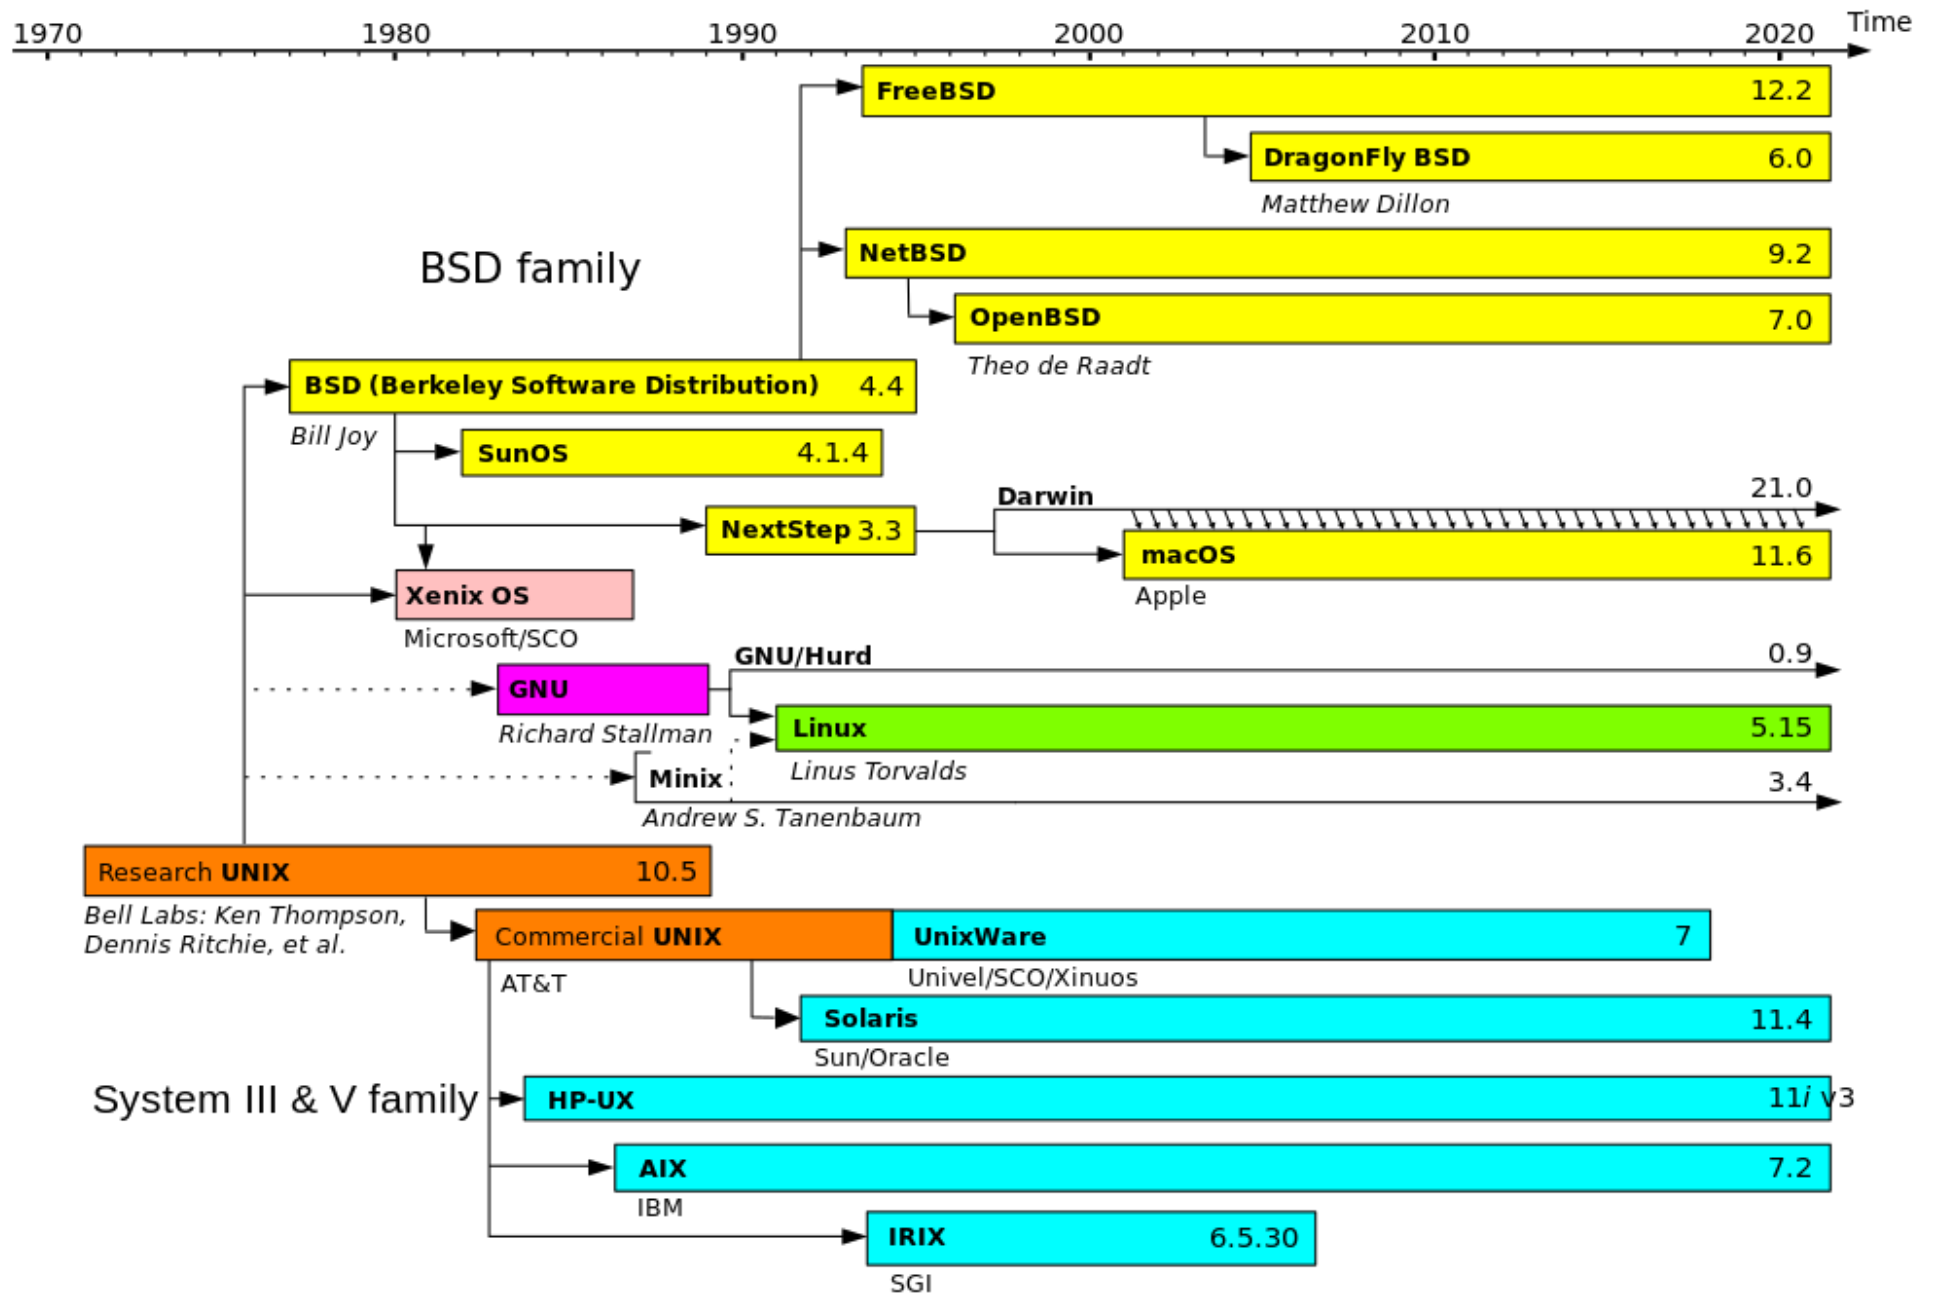
mknod(name, major, minor) Create a device file

fstat(fd) Return info about an open file

link(f1, f2) Create another name (f2) for the file f1

unlink(filename) Remove a file

In many ways xv6 is **very similar** to the operating systems we run today



Evolution of Unix and Unix-like systems



Speakers from the 1984 [Summer](#) USENIX Conference (Salt Lake City, UT)

Backup slides

Pipes

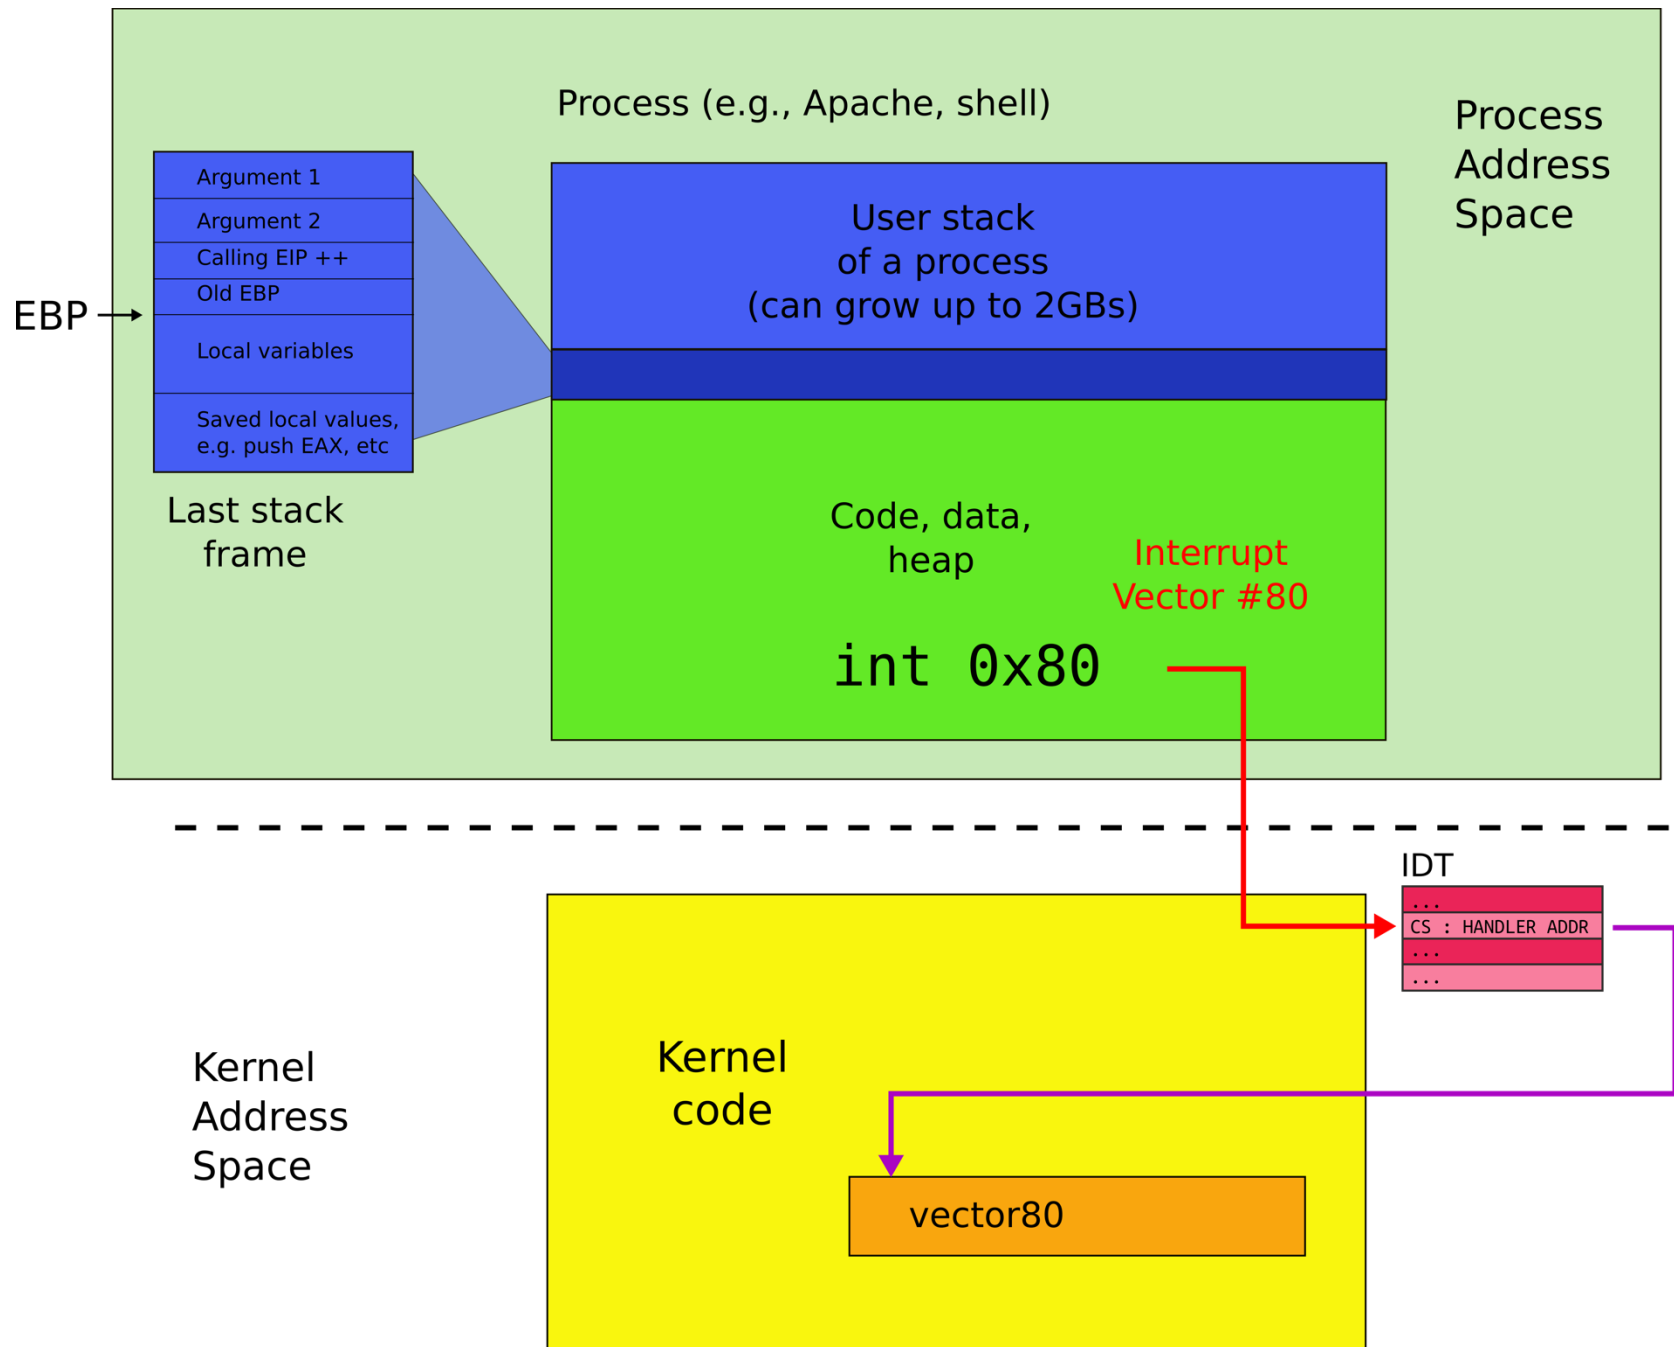
- Shell composes simple utilities into more complex actions with pipes, e.g.

```
grep FORK sh.c | wc -l
```

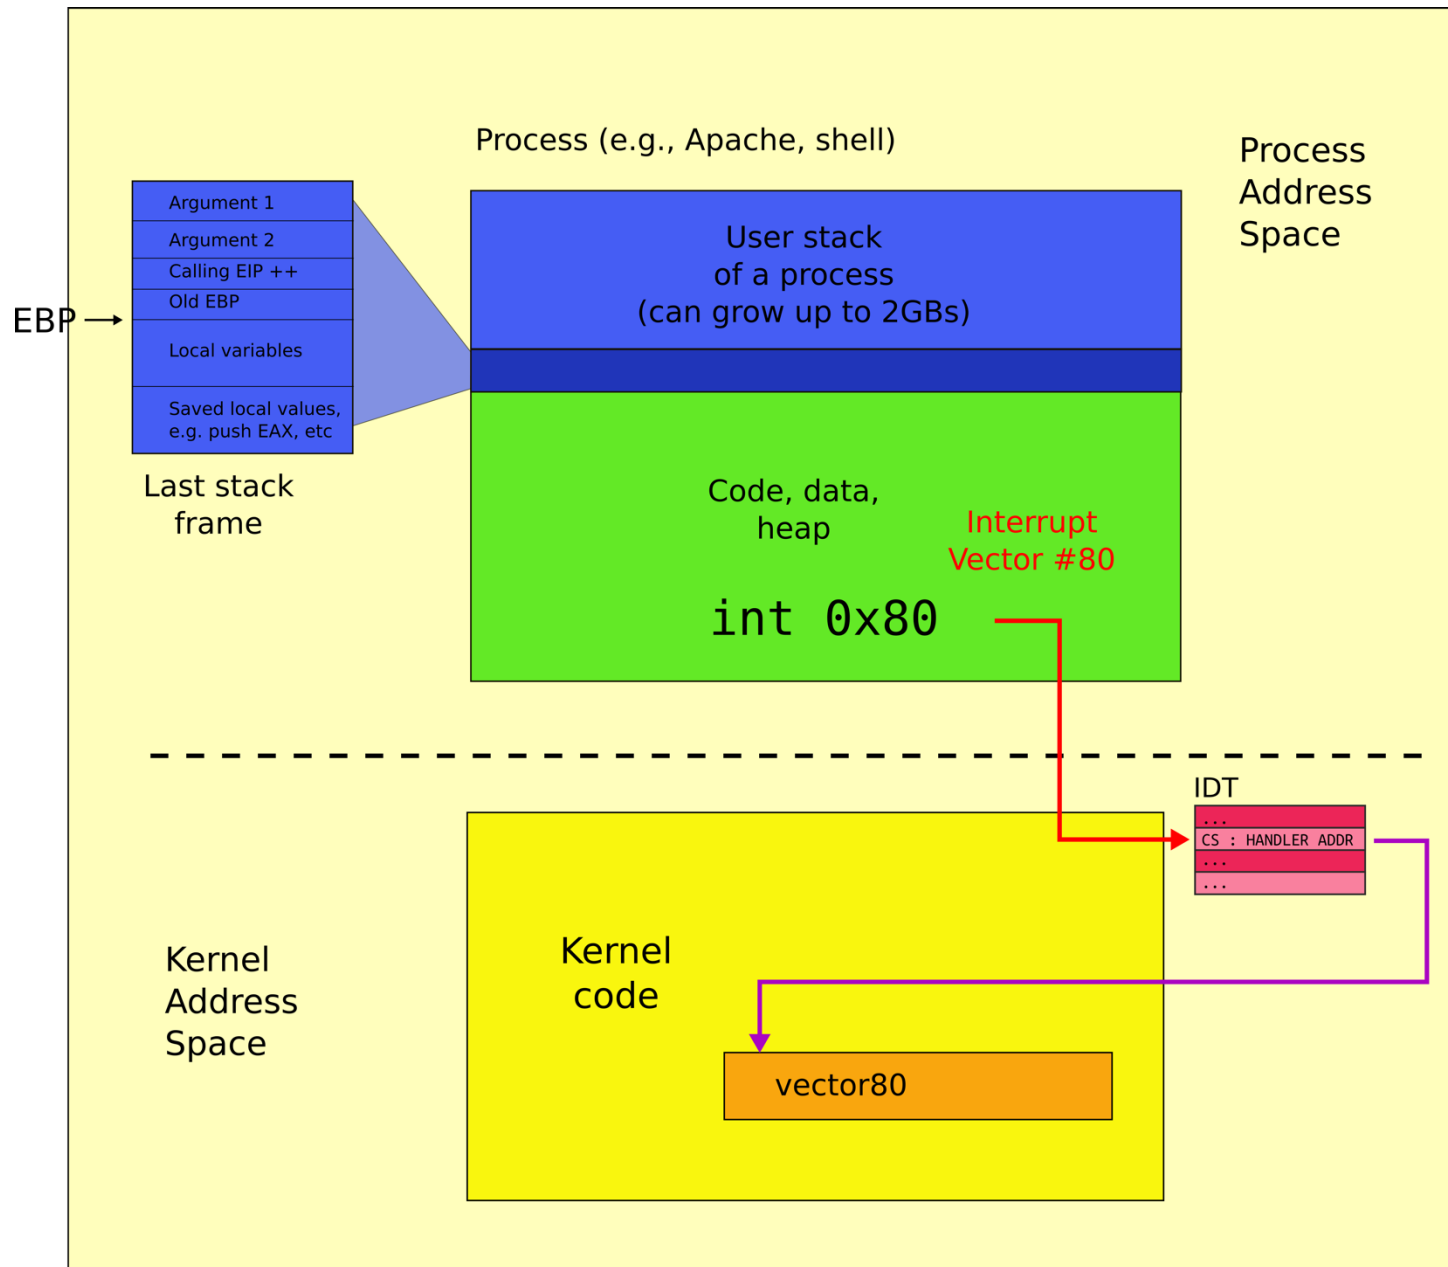
- Create a pipe and connect ends

System call

User address space



Kernel address space



Kernel and user address spaces

