

# cs5460/6460 Operating Systems

## Lecture 4: Function invocations, and calling conventions

Anton Burtsev


January, 2026

# Stack and procedure calls

# What is stack?

# Stack

- It's just a region of memory
  - Pointed by a special register RSP
- You can change RSP
  - Get a new stack



# Why do we need stack?

Stack allows us to invoke functions

# Calling functions

```
// some code...
foo();
// more code..
```


- Stack contains information for **how to return** from a subroutine
- i.e., from **foo()**

- Functions can be called from different places in the program

```
if (a == 0) {
    foo();
    ...
} else {
    foo();
    ...
}
```


# Stack

- Main purpose:
  - Store the return address for the current procedure
  - Caller pushes return address on the stack
  - Callee pops it and jumps



# Stack

- Main purpose:
  - Store the return address for the current procedure
  - Caller pushes return address on the stack
  - Callee pops it and jumps




# Example

```
foo(int a) {  
    if (a == 0)  
        return;  
  
    a--;  
    foo(a);  
    return;  
}  
  
foo(4);
```


# Manipulating stack

- RSP register
- Contains the memory address of the topmost element in the stack
- **PUSH** instruction
  - push 0xBA  
RSP → 0xBA → Stack
- Subtract 8 from RSP
- Insert data on the stack



# Manipulating stack

- **POP instruction**  
`pop RAX`
- Removes data from the stack
- Saves in register or memory
- Adds 8 to RSP



# Calling conventions

# Calling conventions


- Goal: re-entrant programs
- How to pass arguments
  - On the stack?
  - In registers?
- How to return values
  - On the stack?
  - In registers?
- How to save and restore registers across invocations
- Conventions vary between compilers, optimizations, etc.

# Idea 1: Maintain stack as frames

- Each function has a new frame

```
void DrawSquare(...)  
{  
    ...  
    DrawLine(x, y, z);  
}
```

- Use dedicated register **RBP** (frame pointer)
- Points to the base of the frame




# Idea 1: Maintain stack as frames

- Each function has a new frame

```
void DrawSquare(...)  
{  
    ...  
    DrawLine(x, y, z);  
}
```

- Use dedicated register **RBP** (frame pointer)
- Points to the base of the frame




# Idea 1: Maintain stack as frames

- Each function has a new frame

```
void DrawSquare(...)  
{  
    ...  
    DrawLine(x, y, z);  
}
```

- Use dedicated register **RBP** (frame pointer)
- Points to the base of the frame



# Prologue/epilogue

- Each function maintains the frame
- A dedicated register **RBP** is used to keep the frame pointer
- Each function uses **prologue code (blue)**, and **epilogue (green)** to maintain the frame

my\_function:

```
push rbp    ; save original RBP value on stack
mov rbp, rsp ; new RBP = RSP
...         ; function body
pop rbp    ; restore original RBP value
ret
```

# Local variables

# What types of variables do you know?

- Or where these variables are allocated in memory?

# What types of variables do you know?

- **Global** variables
  - Initialized → data section
  - Uninitialized → BSS
- **Dynamic** variables
  - Heap
- **Local** variables
  - Stack

# Global variables

1. #include <stdio.h>
  
2. **char hello[] = "Hello";**
3. int main(int ac, char \*\*av)
4. {
5.   **static char world[] = "world!";**
6.   printf("%s %s\n", hello, world);
7.   return 0;
8. }

# Global variables

1. #include <stdio.h>

2. char hello[] = "Hello";

3. int main(int ac, char \*\*av)

4. {

5. static char world[] = "world!";

6. printf("%s %s\n", hello, world);

7. return 0;

8. }

- Allocated in the data section
- It is split into
  - **Initialized** (non-zero), and **uninitialized** (zero)
  - **Read/write**, and **read-only** data sections

# Global variables

# Dynamic variables (heap)

```
1. #include <stdio.h>
2. #include <string.h>
3. #include <stdlib.h>

4. char hello[] = "Hello";
5. int main(int ac, char **av)
6. {
7.     char world[] = "world!";
8.     char *str = malloc(64);
9.     memcpy(str, "beautiful", 64);
10.    printf("%s %s %s\n", hello, str, world);
11.    return 0;
12. }
```

# Dynamic variables (heap)

```
1. #include <stdio.h>
2. #include <string.h>
3. #include <stdlib.h>

4. char hello[] = "Hello";
5. int main(int ac, char **av)
6. {
7.     char world[] = "world!";
8.     char *str = malloc(64);
9.     memcpy(str, "beautiful", 64);
10.    printf("%s %s %s\n", hello, str, world);
11.    return 0;
12. }
```

- Allocated on the heap
- Special area of memory provided by the OS from where malloc() can allocate memory

# Dynamic variables (heap)

# Local variables

- Local variables

```
1. #include <stdio.h>
2. char hello[] = "Hello";
3. int main(int ac, char **av)
4. {
5.     //static char world[] = "world!";
6.     char world[] = "world!";
7.     printf("%s %s\n", hello, world);
8.     return 0;
9. }
```

# Local variables...

- Each function has private instances of local variables

```
foo(int x) {  
    int a, b, c;  
    ...  
    return;  
}
```

- Function can be called recursively

```
foo(int x) {  
    int a, b, c;  
    a = x + 1;  
    if ( a < 100 )  
        foo(a);  
    return;  
}
```

# How to allocate local variables?

```
foo(int x) {  
    int a, b, c;  
  
    ...  
}
```

# How to allocate local variables?

```
foo(int x) {  
    int a, b, c;  
  
    ...  
}
```

- On the stack!

# Poll Q1: Where do we allocate global variables



<https://pollev.com/cs5460>

# Poll Q2: Where do we allocate dynamic variables



<https://pollev.com/cs5460>

# Allocating local variables

- Stored right after the saved **RBP** value on the stack
- Allocated by subtracting the number of bytes required from **RSP**

\_my\_function:

```
push rbp      ; save original RBP value on stack
mov rbp, rsp  ; new RBP = RSP
sub rsp, LOCAL_BYTES ; locals (round up to keep 16B
                    ; alignment)
...          ; function body
mov rsp, rbp  ; deallocate locals
pop rbp      ; restore original RBP value
ret
```

# Example

```
void my_function() {  
    int a, b, c;  
    ...
```

```
_my_function:  
    push rbp    ; save the value of rbp  
    mov rbp, rsp ; rbp = rsp, set rbp to be top of the stack (rsp)  
    sub rsp, 16  ; move rsp down to allocate space for the  
                 ; local variables on the stack
```

- With frames local variables can be accessed by de-referencing RBP

```
mov [rbp - 4], 10 ; location of variable a  
mov [rbp - 8], 5  ; location of b  
mov [rbp - 12], 2 ; location of c
```

# Example

```
void my_function() {  
    int a, b, c;  
    ...
```

```
_my_function:  
    push rbp    ; save the value of rbp  
    mov rbp, rsp ; rbp = rsp, set rbp to be top of the stack (rsp)  
    sub rsp, 16  ; move rsp down to allocate space for the  
                 ; local variables on the stack
```

- With frames local variables can be accessed by de-referencing RBP

```
mov [rbp - 4], 10 ; location of variable a  
mov [rbp - 8], 5  ; location of b  
mov [rbp - 12], 2 ; location of c
```

# How to pass arguments?

- Possible options:
- In registers
- On the stack

# How to pass arguments?

- Linux x86-64 user-space uses the System V AMD64 ABI
  - Integer/pointer args #1–#6: RDI, RSI, RDX, RCX, R8, R9
  - Floating-point args #1–#8: XMM0–XMM7
- Remaining args: on the stack (8-byte slots)
  - Return values: RAX (and RDX if needed); FP return: XMM0
  - Stack alignment: RSP must be 16-byte aligned before call
  - Red zone: 128 bytes below RSP may be used by leaf functions (user-space)
  - Note: Linux syscalls use a different convention (RAX=syscall number, args RDI/RSI/RDX/R10/R8/R9)

# x86-64 (System V): passing arguments in registers

- Example function

```
void my_function(long x, long y, long z)  
{ ... }
```

- Example invocation

```
my_function(2, 5, 10);
```

- Generated code

```
mov rdi, 2  
mov rsi, 5  
mov rdx, 10  
call my_function
```

# Example stack

```
: :  
| arg8 | [RBP + 24] (8th argument, stack-passed)  
| arg7 | [RBP + 16] (7th argument, stack-passed)  
| RA  | [RBP + 8] (return address)  
| FP  | [RBP]    (old RBP value) ← RBP points here  
| ... | [RBP - 8] (1st local variable)  
:  
:  
| ... | [RBP - X] (RSP - the current stack pointer)
```

args 1–6 are in registers:

RDI, RSI, RDX, RCX, R8, R9

# Example: caller side code

```
long callee(long, long, long);

long caller(void)
{
    long ret;

    ret = callee(1, 2, 3);
    ret += 5;
    return ret;
}
```

```
caller:
; manage own stack frame
push rbp
mov rbp, rsp
; set call arguments (registers)
mov rdi, 1
mov rsi, 2
mov rdx, 3
; call subroutine 'callee'
call callee
; no stack cleanup needed
; (args in regs)
; stack is 16B-aligned before call
; use subroutine result
add rax, 5
; restore old call frame
pop rbp
; return
ret
```

# Example: caller side code

```
long callee(long, long, long);
```

```
long caller(void)
```

```
{
```

```
    long ret;
```

```
    ret = callee(1, 2, 3);
```

```
    ret += 5;
```

```
    return ret;
```

```
}
```

caller:

; manage own stack frame

push rbp

mov rbp, rsp

; set call arguments (registers)

mov rdi, 1

mov rsi, 2

mov rdx, 3

; call subroutine 'callee'

call callee

; no stack cleanup needed

; (args in regs)

; stack is 16B-aligned before call

; use subroutine result

add rax, 5

; restore old call frame

pop rbp

; return

ret

# Example: caller side code

```
long callee(long, long, long);
```

```
long caller(void)
```

```
{
```

```
    long ret;
```

```
    ret = callee(1, 2, 3);
```

```
    ret += 5;
```

```
    return ret;
```

```
}
```

caller:

```
; manage own stack frame  
push rbp  
mov rbp, rsp  
; set call arguments (registers)  
mov rdi, 1  
mov rsi, 2  
mov rdx, 3  
; call subroutine 'callee'  
call callee  
; no stack cleanup needed  
; (args in regs)  
; stack is 16B-aligned before call  
; use subroutine result  
add rax, 5  
; restore old call frame  
pop rbp  
; return  
ret
```

# Example: caller side code

```
long callee(long, long, long);

long caller(void)
{
    long ret;

    ret = callee(1, 2, 3);
    ret += 5;
    return ret;
}
```

caller:

```
; manage own stack frame
push rbp
mov rbp, rsp
; set call arguments (registers)
mov rdi, 1
mov rsi, 2
mov rdx, 3
; call subroutine 'callee'
call callee
; no stack cleanup needed
; (args in regs)
; stack is 16B-aligned before call
; use subroutine result
add rax, 5
; restore old call frame
pop rbp
; return
ret
```

# Example: caller side code

```
long callee(long, long, long);

long caller(void)
{
    long ret;

    ret = callee(1, 2, 3);
    ret += 5;
    return ret;
}
```

caller:

```
; manage own stack frame
push rbp
mov rbp, rsp
; set call arguments (registers)
mov rdi, 1
mov rsi, 2
mov rdx, 3
; call subroutine 'callee'
call callee
; no stack cleanup needed
; (args in regs)
; stack is 16B-aligned before call
; use subroutine result
add rax, 5
; restore old call frame
pop rbp
; return
ret
```

# Example: caller side code

```
long callee(long, long, long);

long caller(void)
{
    long ret;

    ret = callee(1, 2, 3);
    ret += 5;
    return ret;
}
```

caller:

```
; manage own stack frame
push rbp
mov rbp, rsp
; set call arguments (registers)
mov rdi, 1
mov rsi, 2
mov rdx, 3
; call subroutine 'callee'
call callee
; no stack cleanup needed
; (args in regs)
; stack is 16B-aligned before call
; use subroutine result
add rax, 5
; restore old call frame
pop rbp
; return
ret
```

# Example: callee side code

```
void my_function(long x, long y, long z)
{
    long a, b, c;
    ...
    return;
}
```

```
my_function:
    push rbp
    mov rbp, rsp
    sub rsp, 32 ; allocate locals + padding (keep 16B alignment)
        ; x in RDI, y in RSI, z in RDX (System V AMD64)
        ; a=[rbp-8], b=[rbp-16], c=[rbp-24]
```


```
    mov rsp, rbp ; deallocate local variables
    pop rbp
    ret
```

# Back to stack frames, so why do we need them?

- ... They are not strictly required
- GCC compiler option `-fomit-frame-pointer` can disable them

Don't keep the frame pointer in a register for functions that don't need one. This avoids the instructions to save, set up and restore frame pointers; it also makes an extra register available in many functions. **It also makes debugging impossible on some machines.**

# Referencing args without frames



- In x86-64 SysV, most arguments arrive in registers
- Stack-passed args ( $\text{arg}_7+$ ) start at [**RSP + 8**] on entry (above return address)
- As you push/pop or allocate locals, **RSP changes** → offsets change too

- Debugging becomes hard
  - As RSP changes one has to manually keep track where local variables are relative to RSP (`RSP + 4` or `+8`)
- **Compiler can easily do this and generate correct code!**
- **But it's hard for a human**
  - It's hard to unwind the stack in case of a crash
  - To print out a backtrace

# And you only save...

- A stack frame typically costs a couple instructions + one register (RBP as frame pointer)
  - x86-64 has **16 general-purpose registers**, so keeping RBP as a frame pointer is usually **affordable**
  - Optimizing compilers often omit the frame pointer by default (-fomit-frame-pointer)
  - GCC/Clang at -O1/-O2/... typically enable that option
  - For easier backtraces/profiling: compile with -fno-omit-frame-pointer

# Relevant part of the GCC manual

## 3.10 Options That Control Optimization

[https://gcc.gnu.org/onlinedocs/gcc/Optimize-  
Options.html](https://gcc.gnu.org/onlinedocs/gcc/Optimize-Options.html)

# Poll

- In the calling convention we just discussed how do we pass the first argument to the function?



<https://pollev.com/cs5460>

# Saving and restoring registers

# Saving register state across invocations

- Processor doesn't save registers
- General purpose, segment, flags
- Again, a calling convention is needed
- Agreement on what gets saved by the callee and the caller

# Saving register state across invocations

- System V AMD64 (Linux): **caller-saved** (volatile): RAX, RCX, RDX, RSI, RDI, R8–R11 (+ XMM0–XMM15)
  - Callee-saved (non-volatile): RBX, RBP, R12–R15 (and RSP must be restored)
- **Rule:** if the callee uses a callee-saved register, it must restore it before returning
  - Caller must assume caller-saved regs are clobbered across a call

- In general there are multiple calling conventions / ABIs
- On Linux x86-64 user-space the standard is the System V AMD64 ABI
- 32-bit code often used cdecl/stdcall/fastcall (different rules)
- Docs:  
[https://en.wikipedia.org/wiki/X86\\_calling\\_conventions#System\\_V\\_AMD64\\_ABI](https://en.wikipedia.org/wiki/X86_calling_conventions#System_V_AMD64_ABI)
- Be careful when mixing C/assembly, calling into libraries, or crossing ABIs

# References

- [https://en.wikibooks.org/wiki/X86\\_Disassembly/Functions\\_and\\_Stack\\_Frames](https://en.wikibooks.org/wiki/X86_Disassembly/Functions_and_Stack_Frames)
- [https://en.wikipedia.org/wiki/Calling\\_convention](https://en.wikipedia.org/wiki/Calling_convention)
- [https://en.wikipedia.org/wiki/X86\\_calling\\_conventions](https://en.wikipedia.org/wiki/X86_calling_conventions)
- <http://stackoverflow.com/questions/14666665/trying-to-understand-gcc-option-fomit-frame-pointer>
- [https://en.wikipedia.org/wiki/X86\\_calling\\_conventions#System\\_V\\_AMD64\\_ABI](https://en.wikipedia.org/wiki/X86_calling_conventions#System_V_AMD64_ABI)

# Thank you!

