

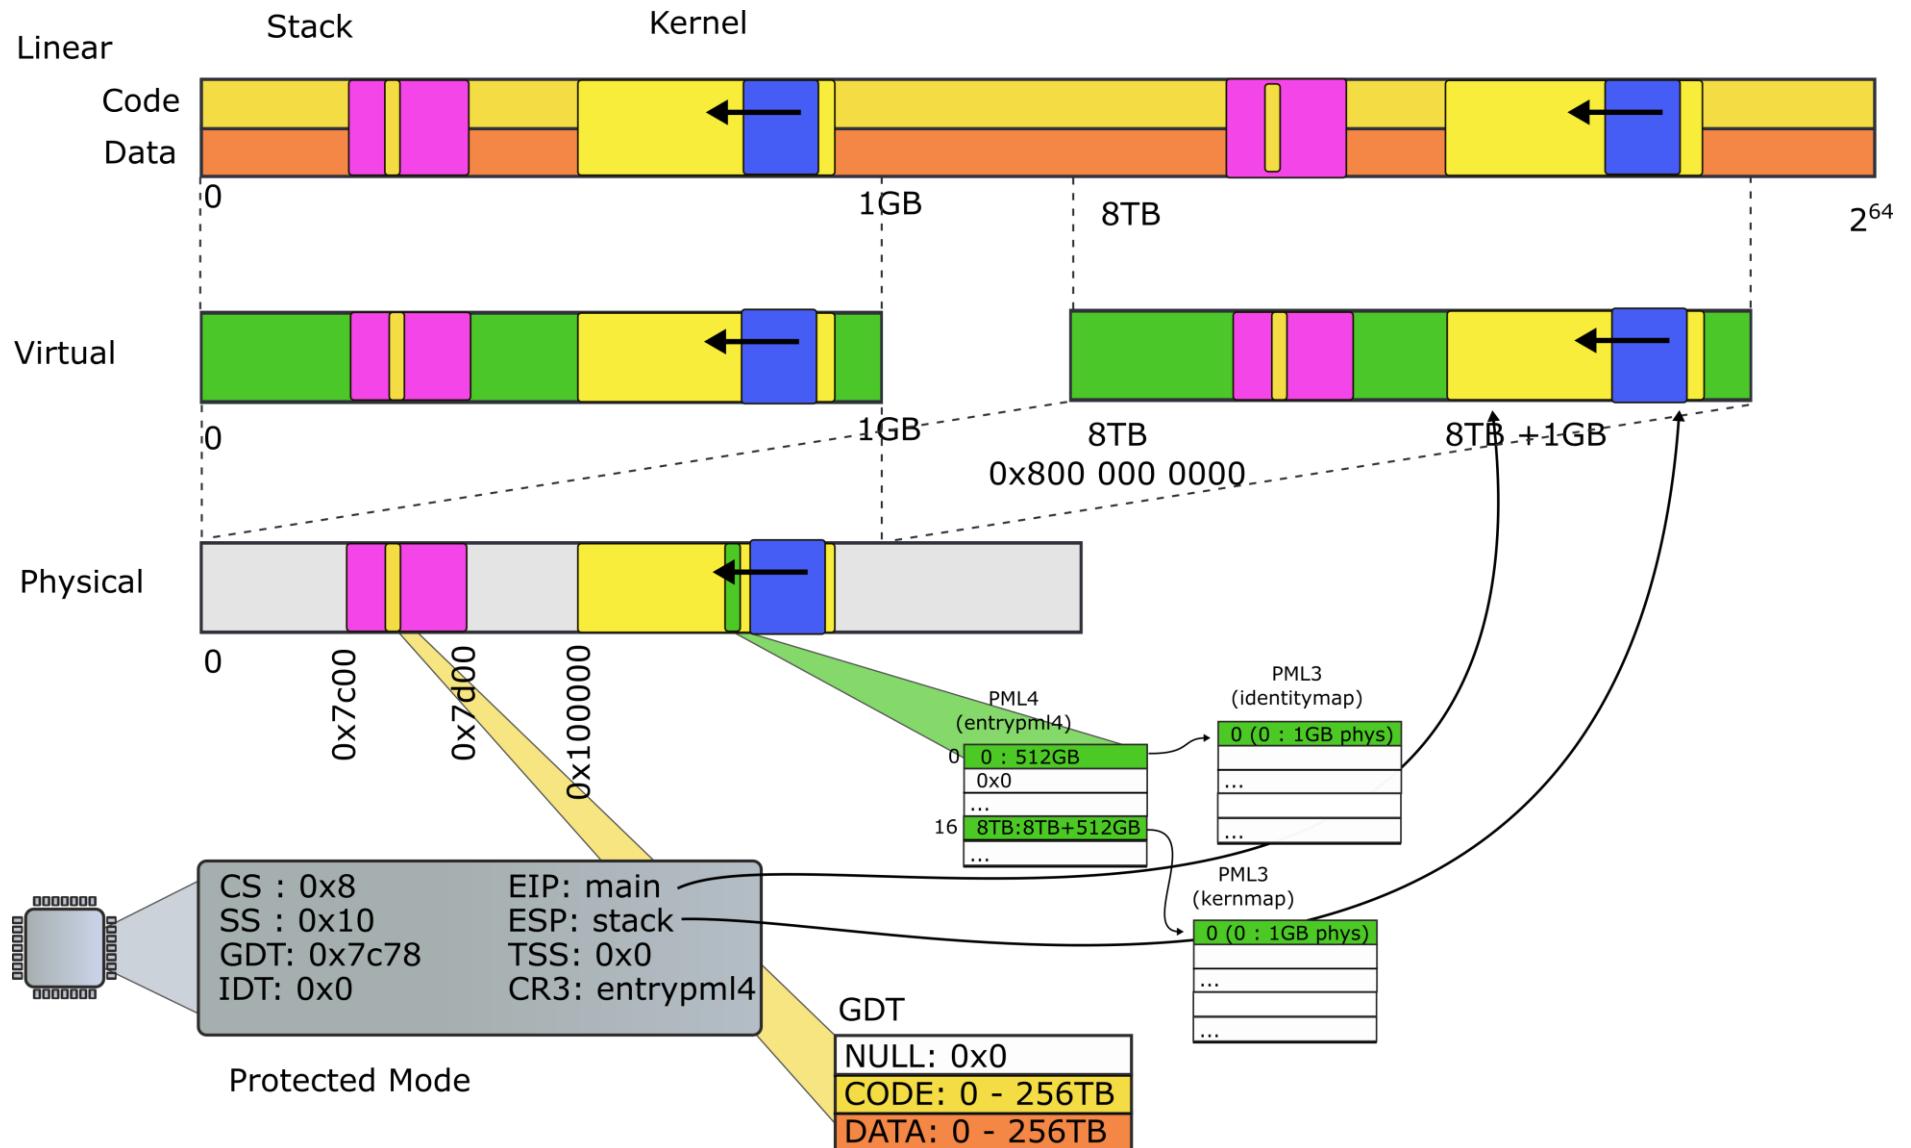
# cs5460/6460: Operating Systems

## Lecture 08: System Init (Kernel Memory Allocator and Page Table)

Anton Burtsev

February, 2026

# State of the system after boot



# Running in main()

1313 // Bootstrap processor starts running C code here.

1314 // Allocate a real stack and switch to it, first

1315 // doing some setup required for memory allocator to work.

1316 int

1317 main(void)

1318 {

1319   kinit1(end, P2V(4\*1024\*1024)); // phys page allocator

1320   kvmalloc(); // kernel page table

1321   mpinit(); // detect other processors

1322   lapicinit(); // interrupt controller

1323   seginit(); // segment descriptors


1324   cprintf("\ncpu%d: starting xv6\n\n", cpunum());

...

1340 }

What's next?

# We want to run multiple processes



But what is a process?


# A couple of requirements

- Each process is a collection of resources
- Memory
  - E.g., text, stack, heap
- In-kernel state
  - E.g., open file descriptors, network sockets (connections)

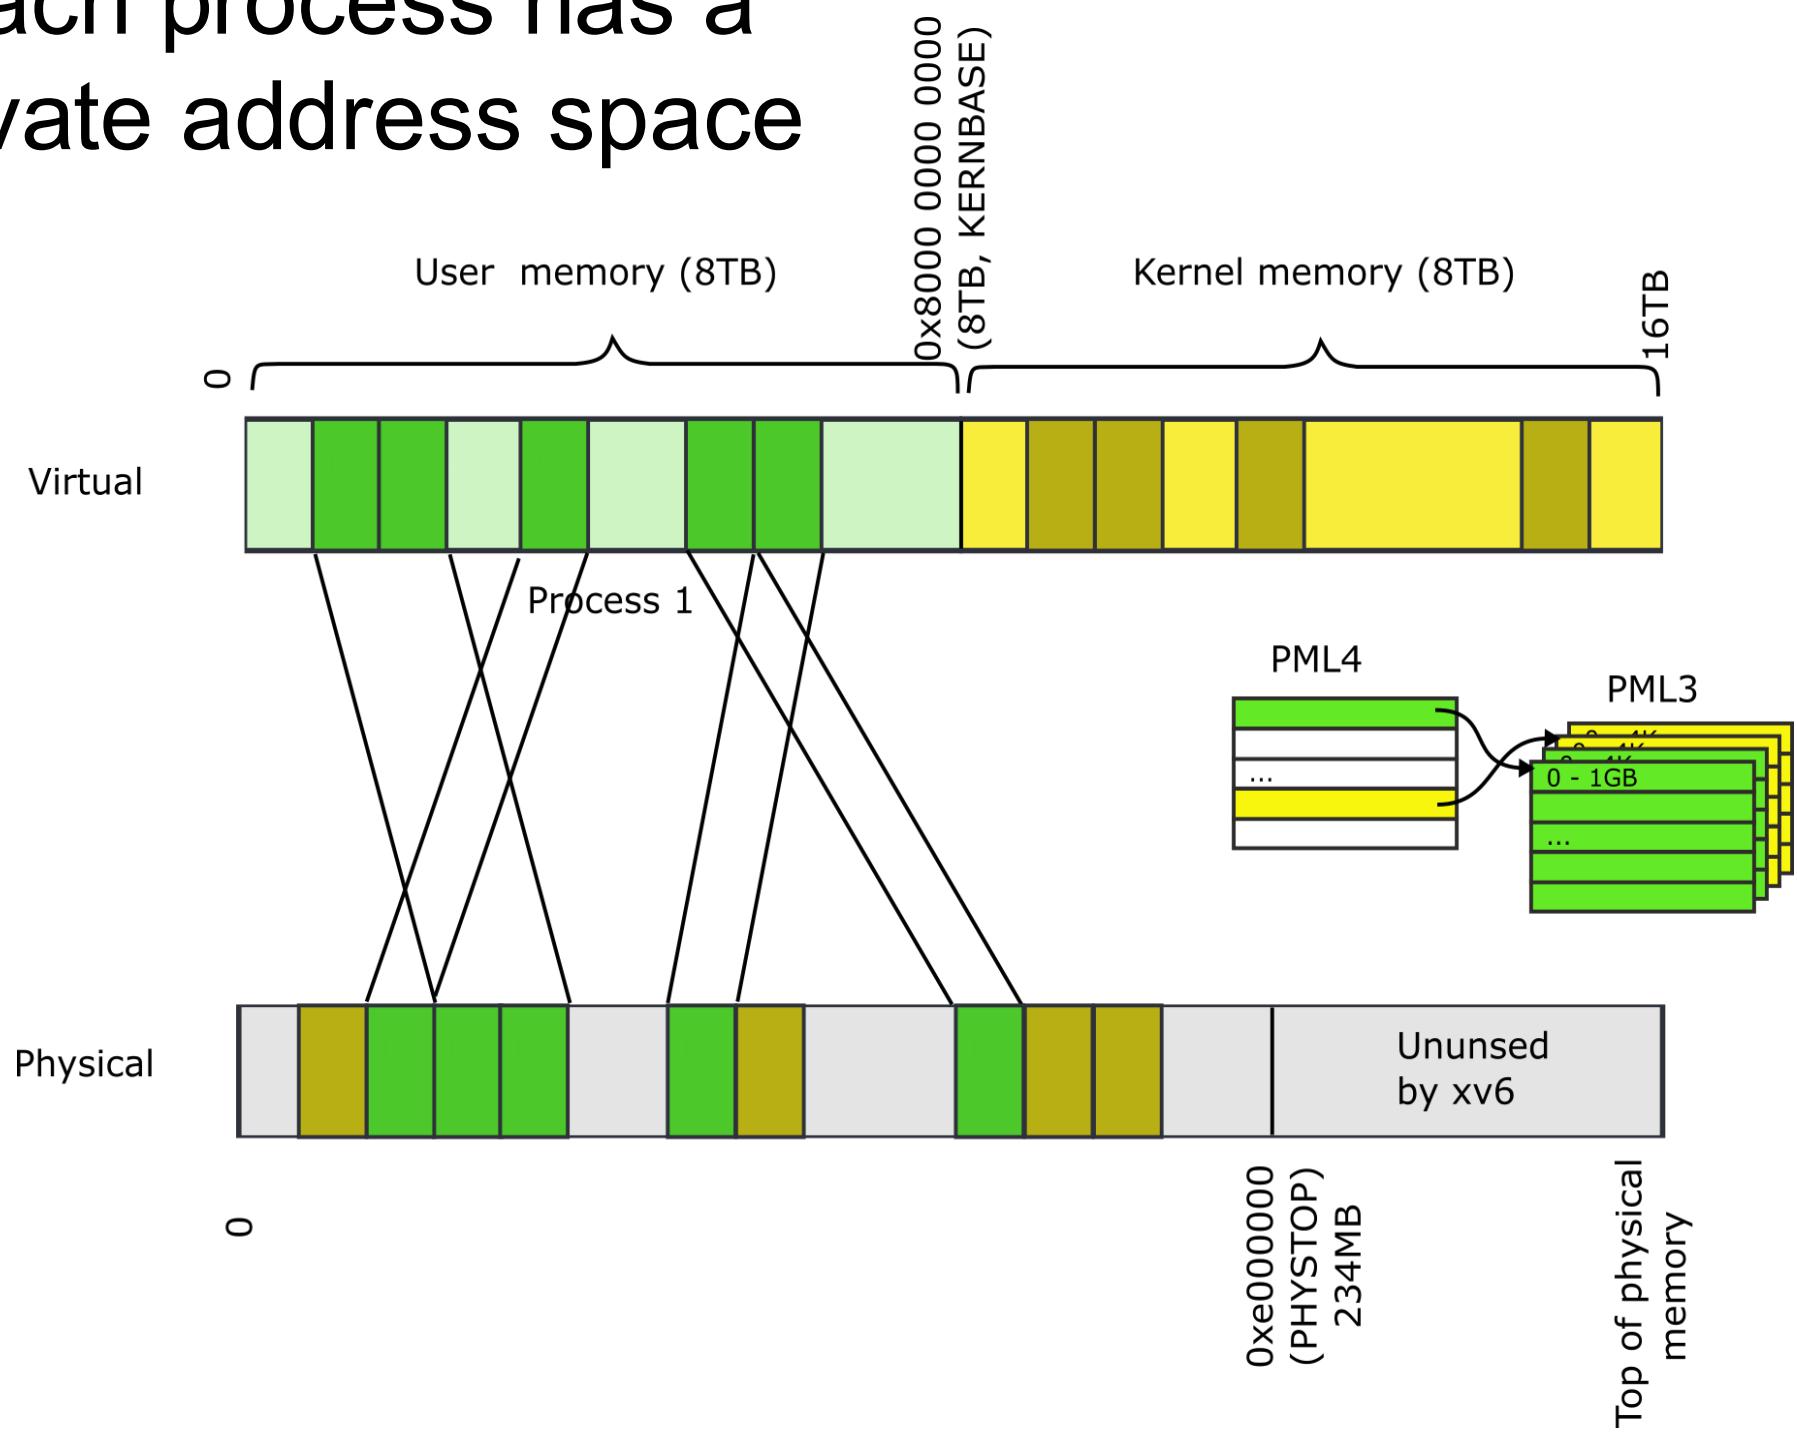
# A couple of requirements

- Each process is a collection of resources
- Memory
  - E.g., text, stack, heap
- In-kernel state
  - E.g., open file descriptors, network sockets (connections)
- Processes are **isolated** from each other
- Processes **don't trust** each other
  - Individual users, some privileged
- Can't interfere with other processes
- Can't change kernel (to affect other processes)

# Each process will have 8TB private address space

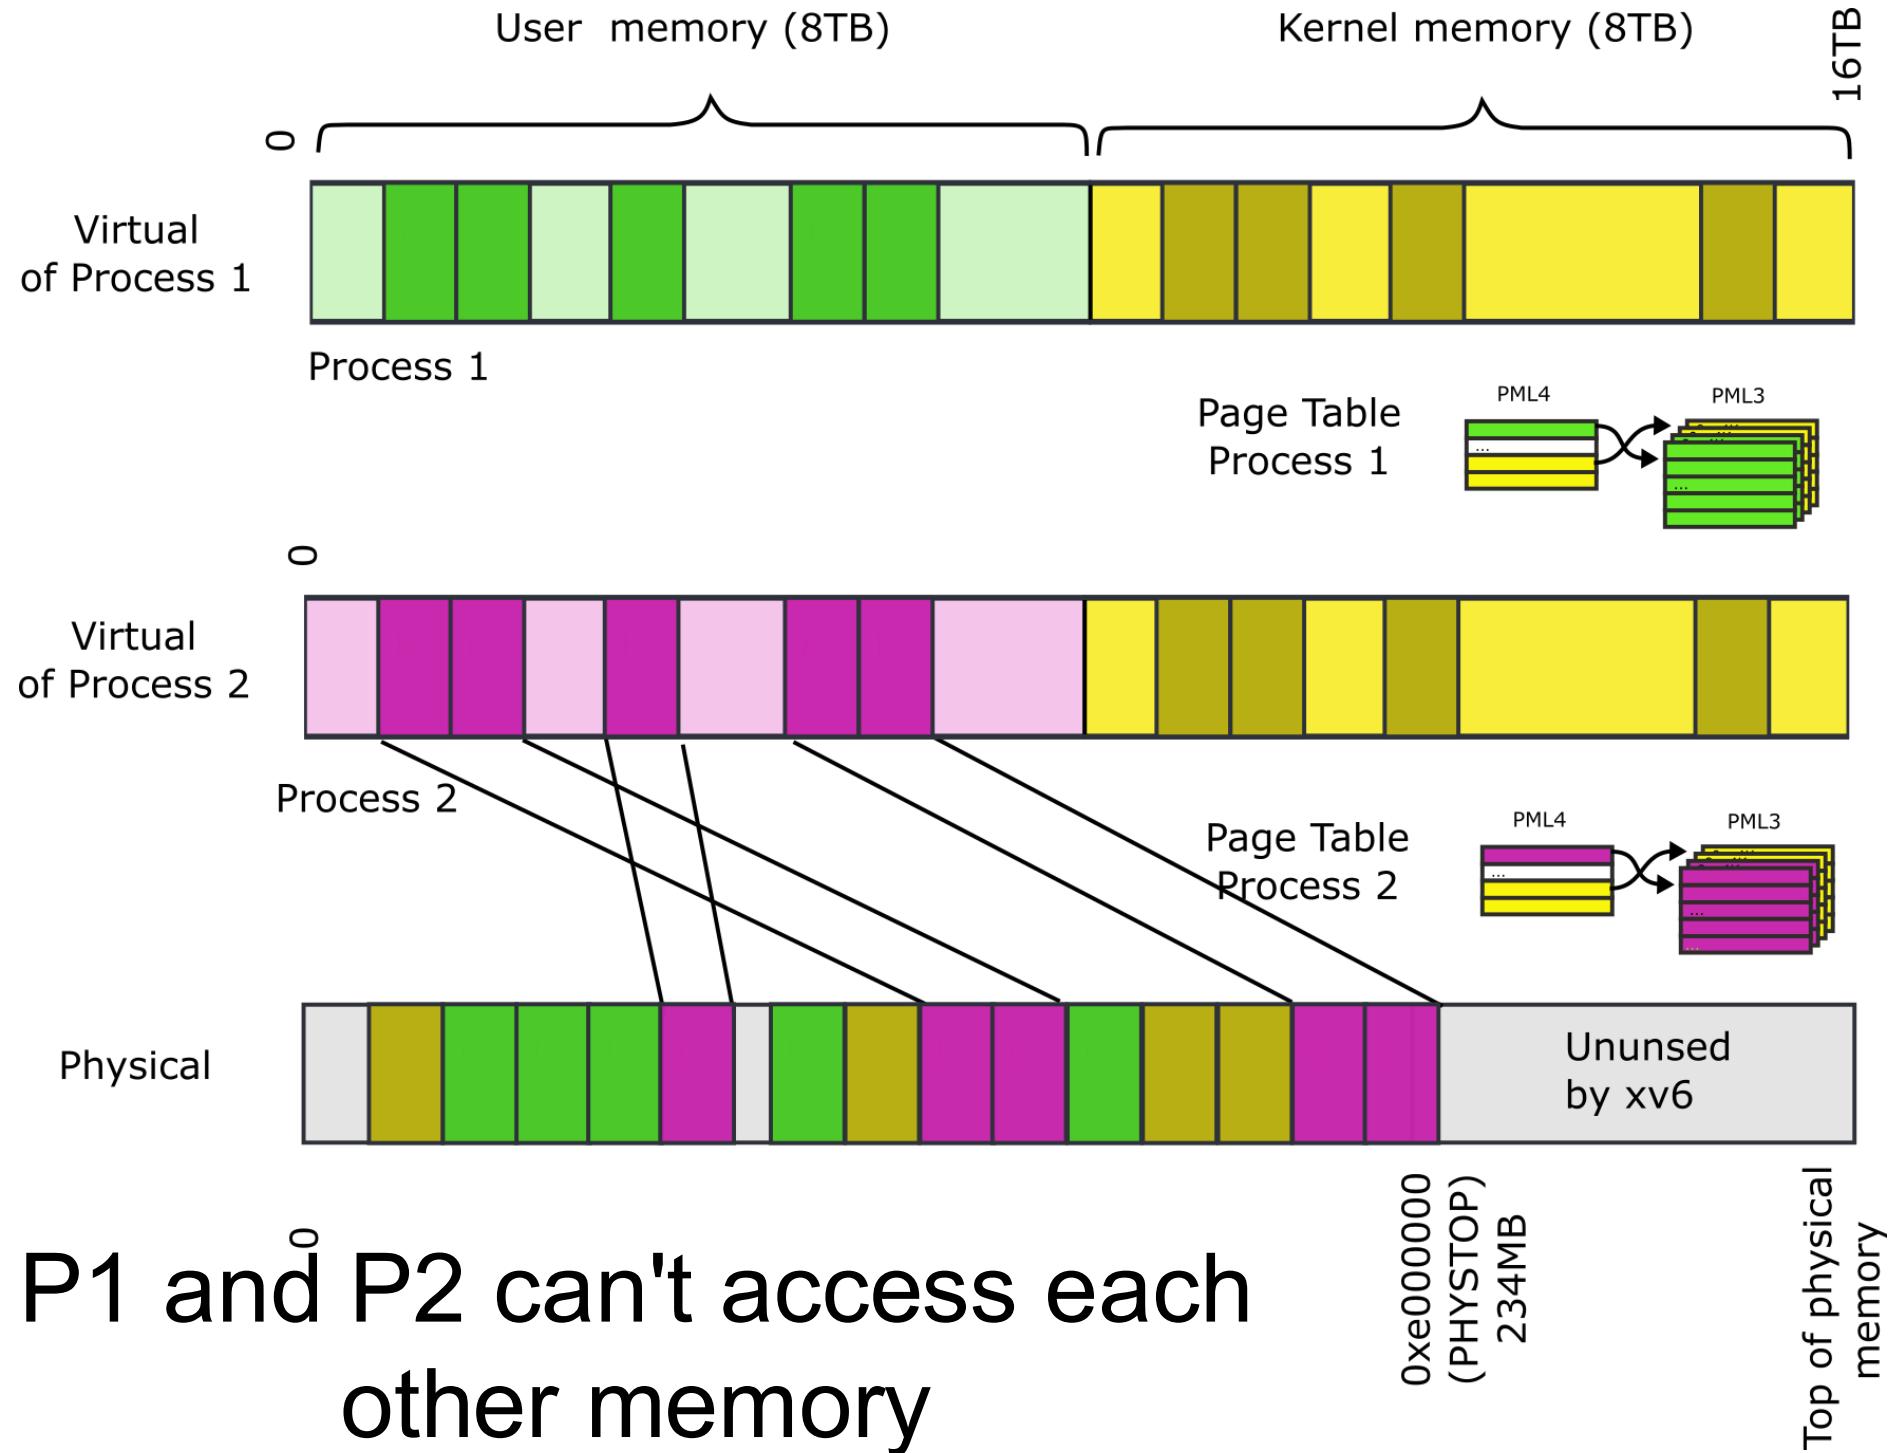


# Each process has a private address space

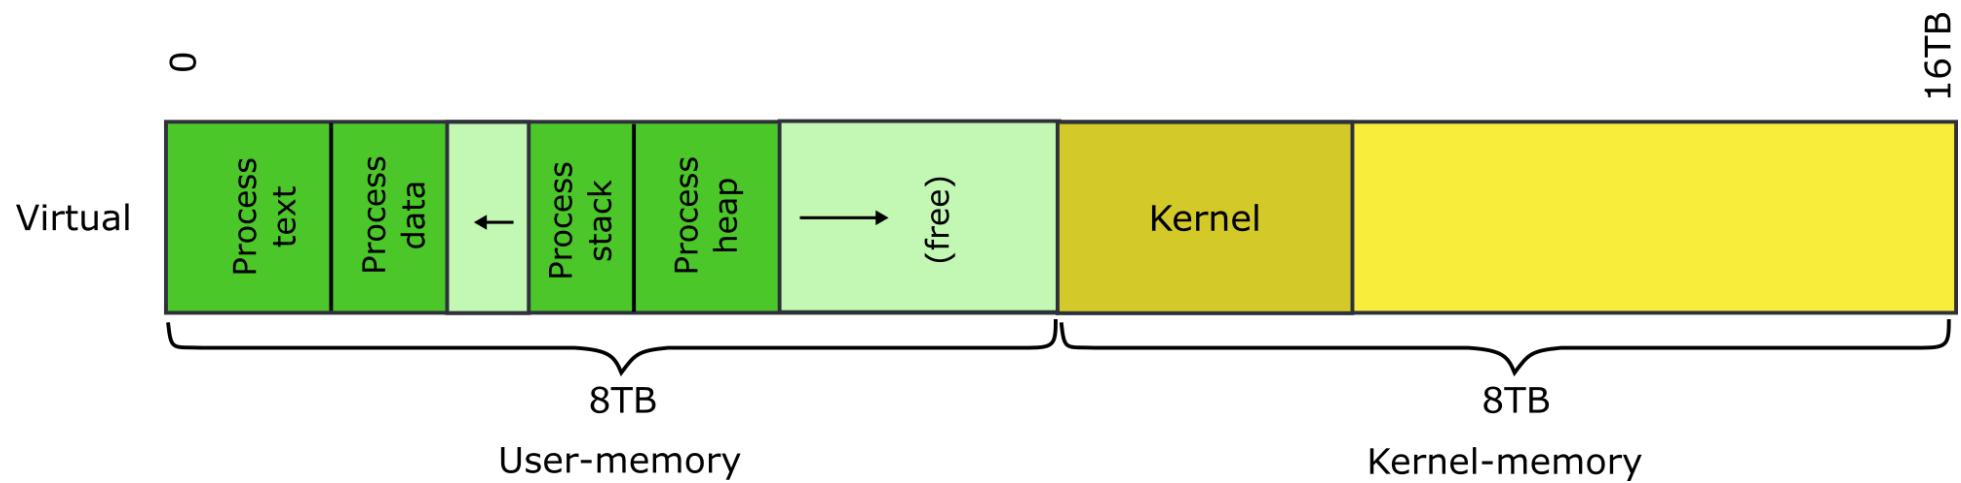


# Each process maps the kernel


- It's not strictly required
- But convenient for system calls
- No need to change the page table when process enters the kernel with a system call
- **Things are much faster!**



# Our goal: split address space



# Memory after boot



# Outline

- Create the kernel address space
- Create kernel memory allocator
- Allocate memory for page tables
  - Page table directory and page table

# Kernel memory allocator

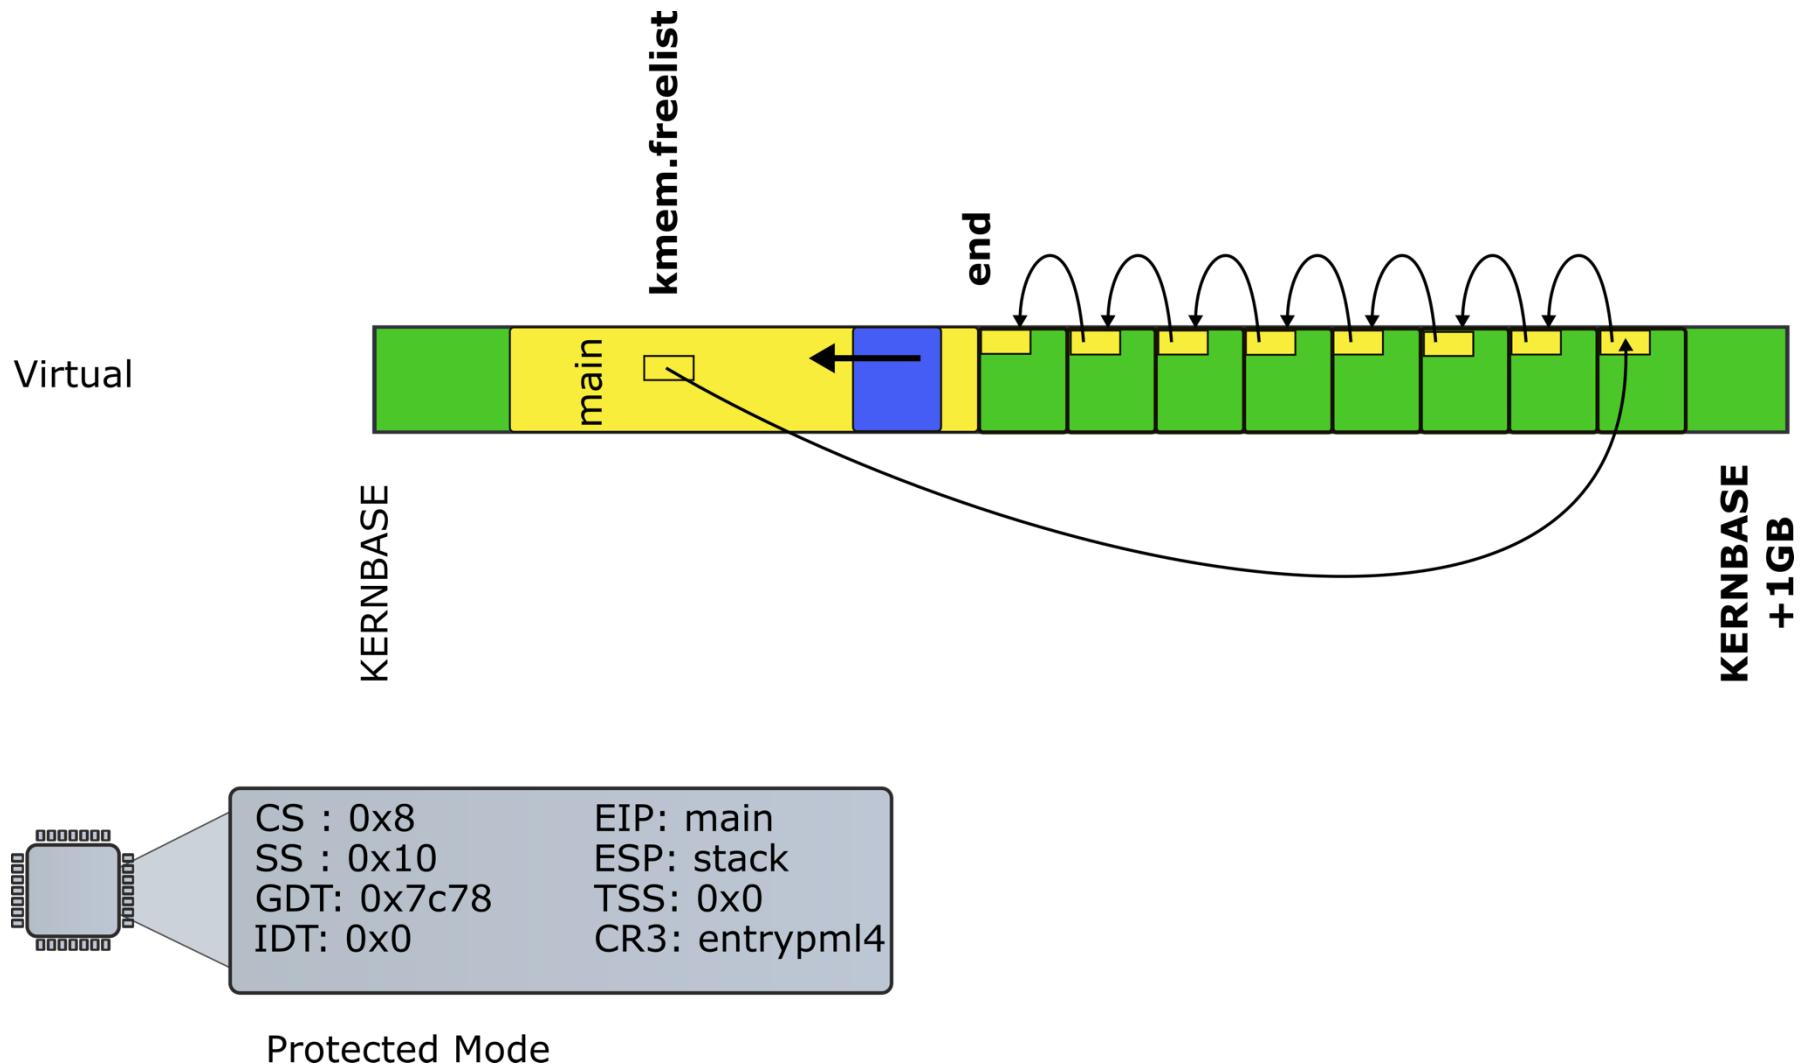
- Kernel needs normal 4-level, 4KB page table
- Right now we have
  - One (statically allocated) page table
  - That has only two 1GB entries
- 4KB page table is a better choice
  - Xv6 processes are small
  - Wasting 1GB or 2MB on a program that fits into 1KB is absurd
- But to create page tables we need memory
- Where can it come from?

# Simple memory allocator

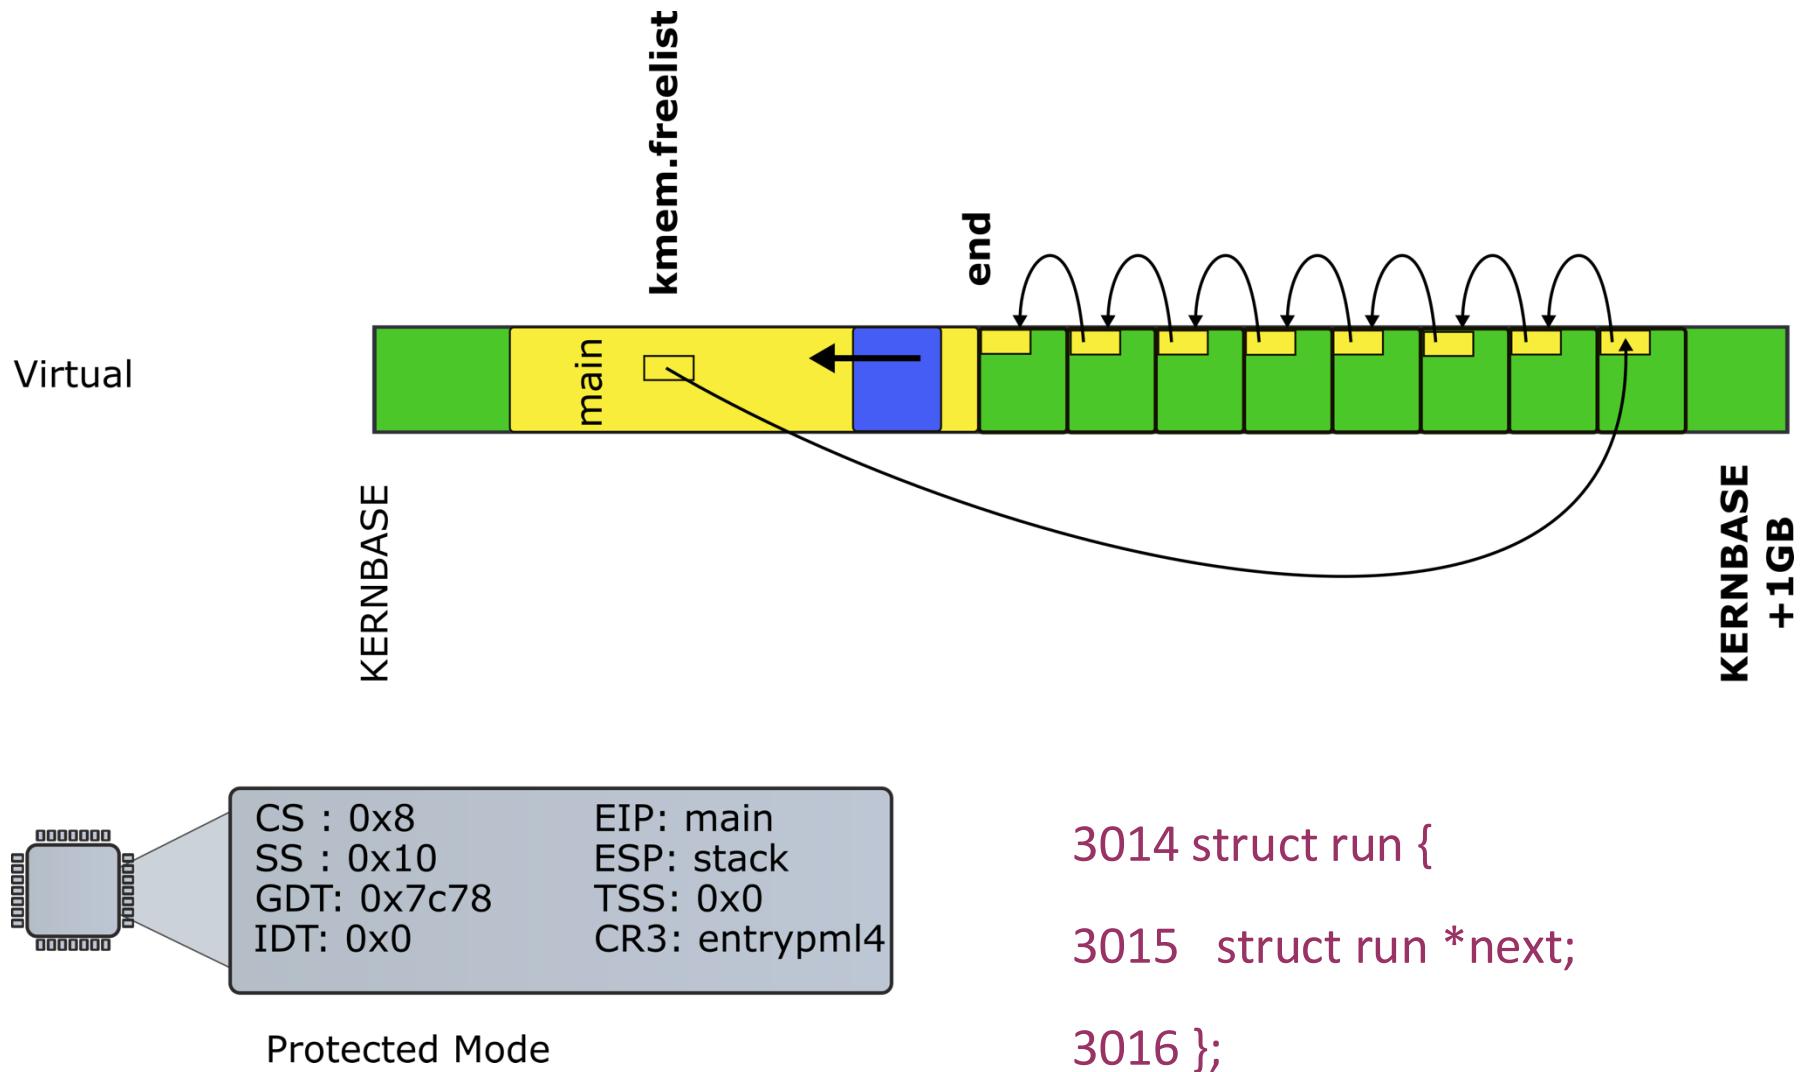
- Goal:
  - `alloc()` and `free()`
  - To allocate page tables, stacks, data structures, etc.

What can it look like?

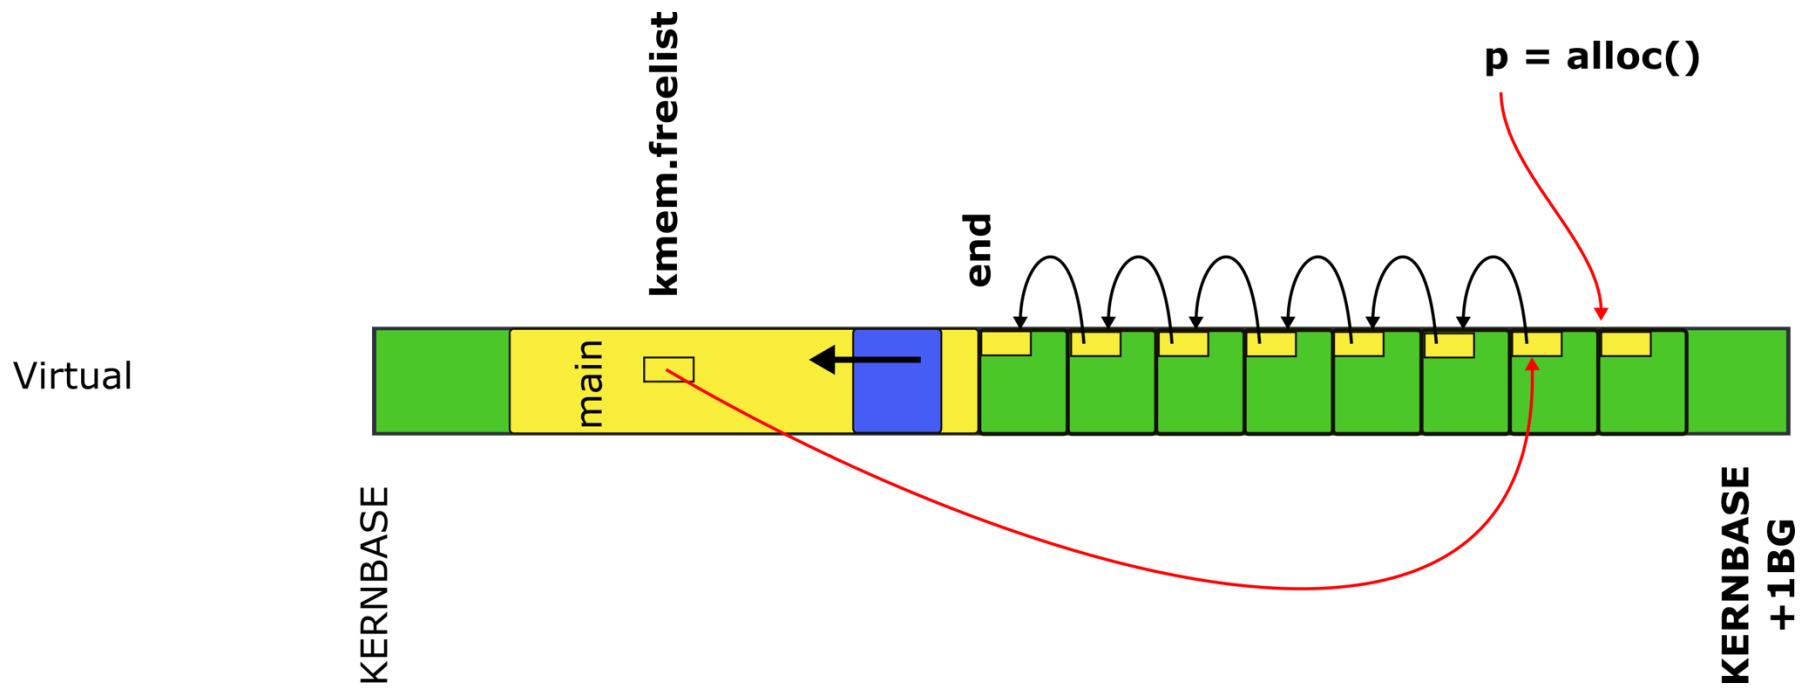
# Page allocator



# Page allocator



# Page allocator

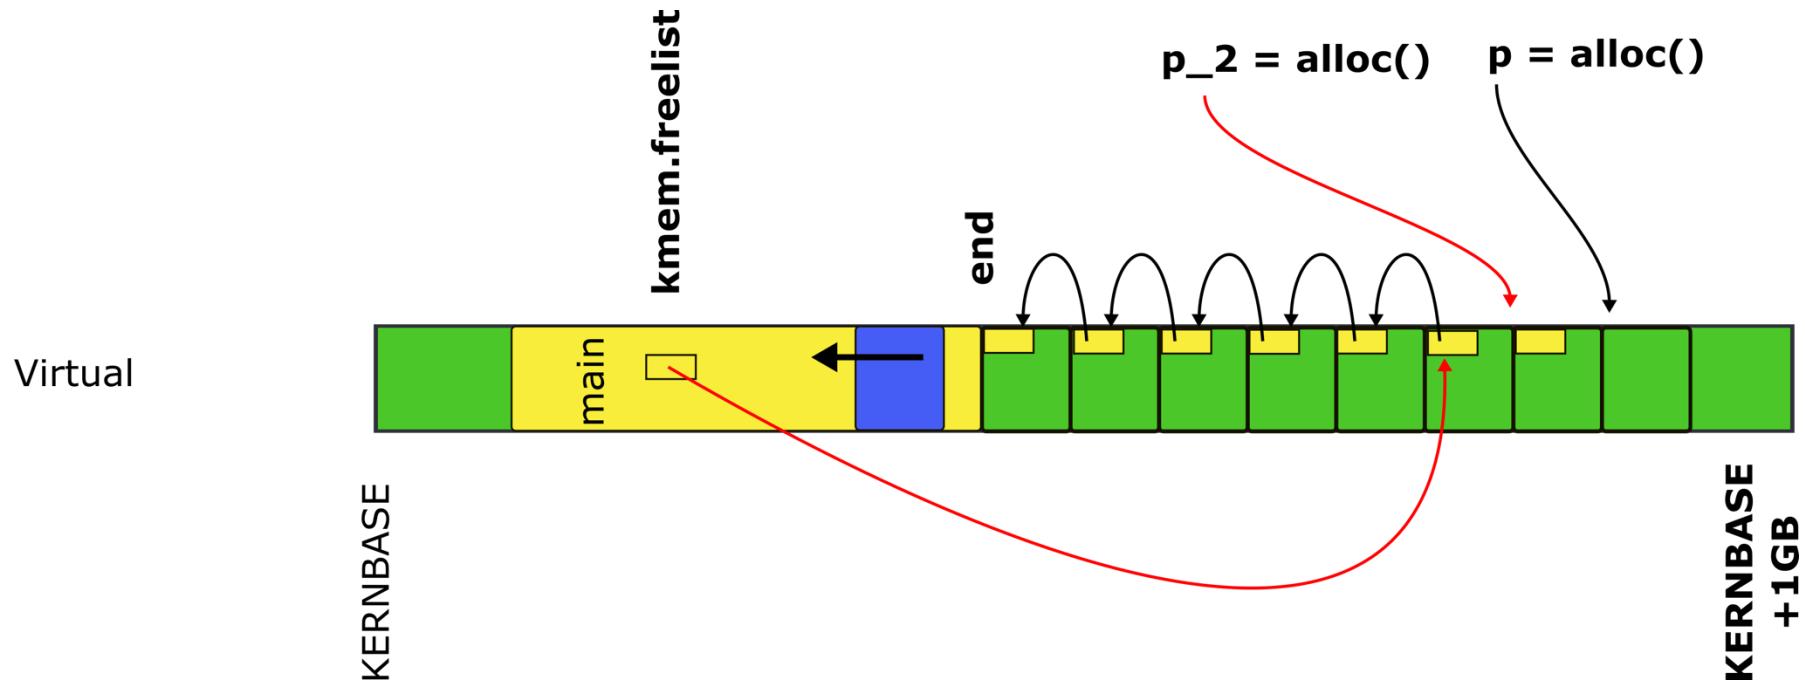


3014 struct run {

3015 struct run \*next;

3016};

# Page allocator

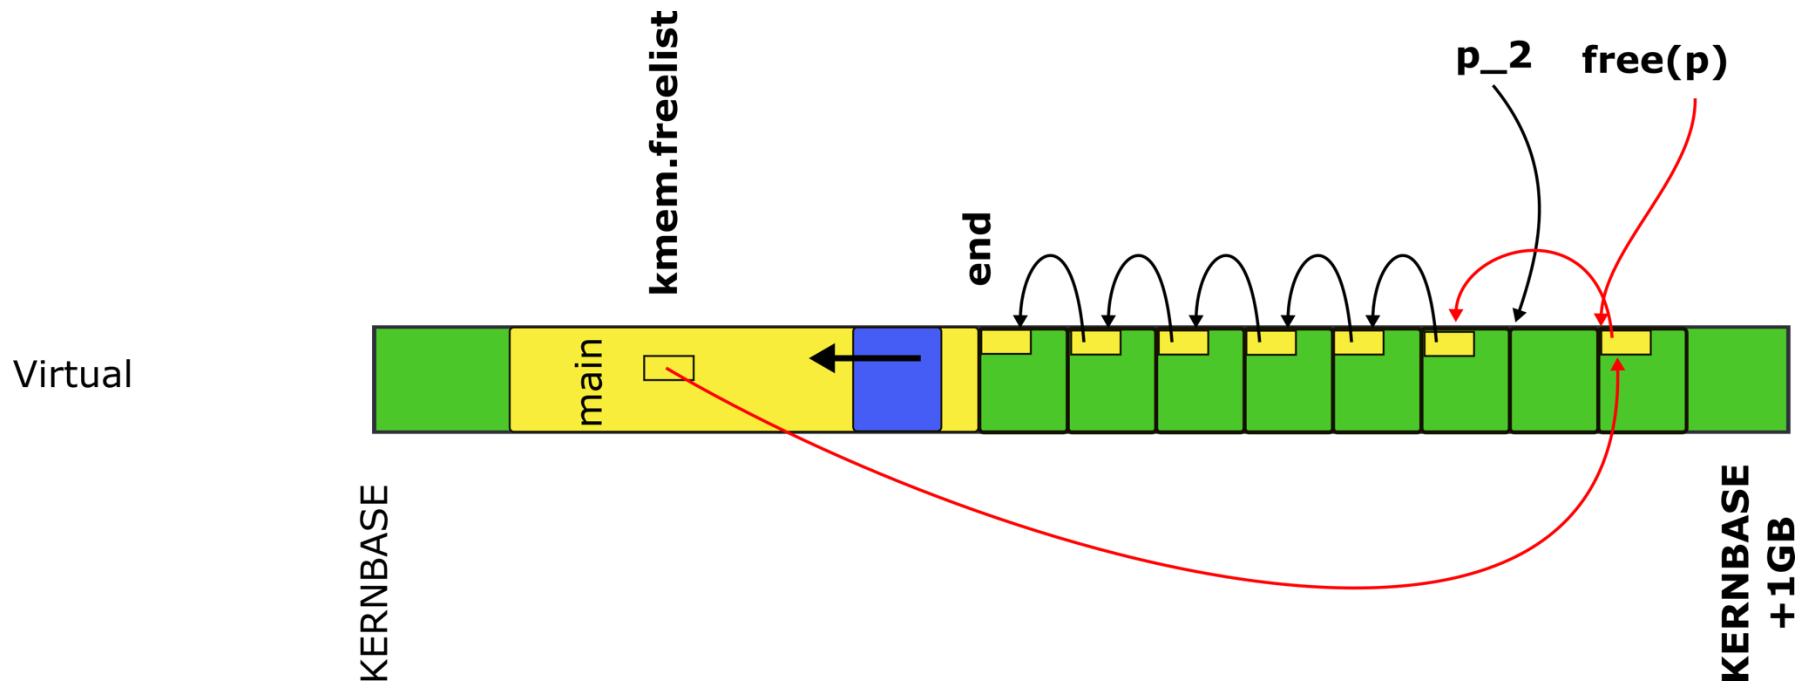


3014 struct run {

3015 struct run \*next;

3016};

# Page allocator



3014 struct run {

3015 struct run \*next;

3016};

# kalloc() - kernel allocator

3087 char\*

3088 kalloc(void)

3089 {

3080 struct run \*r;

...

3094 r = kmem.freelist;


3095 if(r)


3096 kmem.freelist = r->next;

...

3099 return (char\*)r;

3099 }





```
3065 kfree(char *v)
```

```
3066 {
```

```
3067 struct run *r;
```

```
...
```

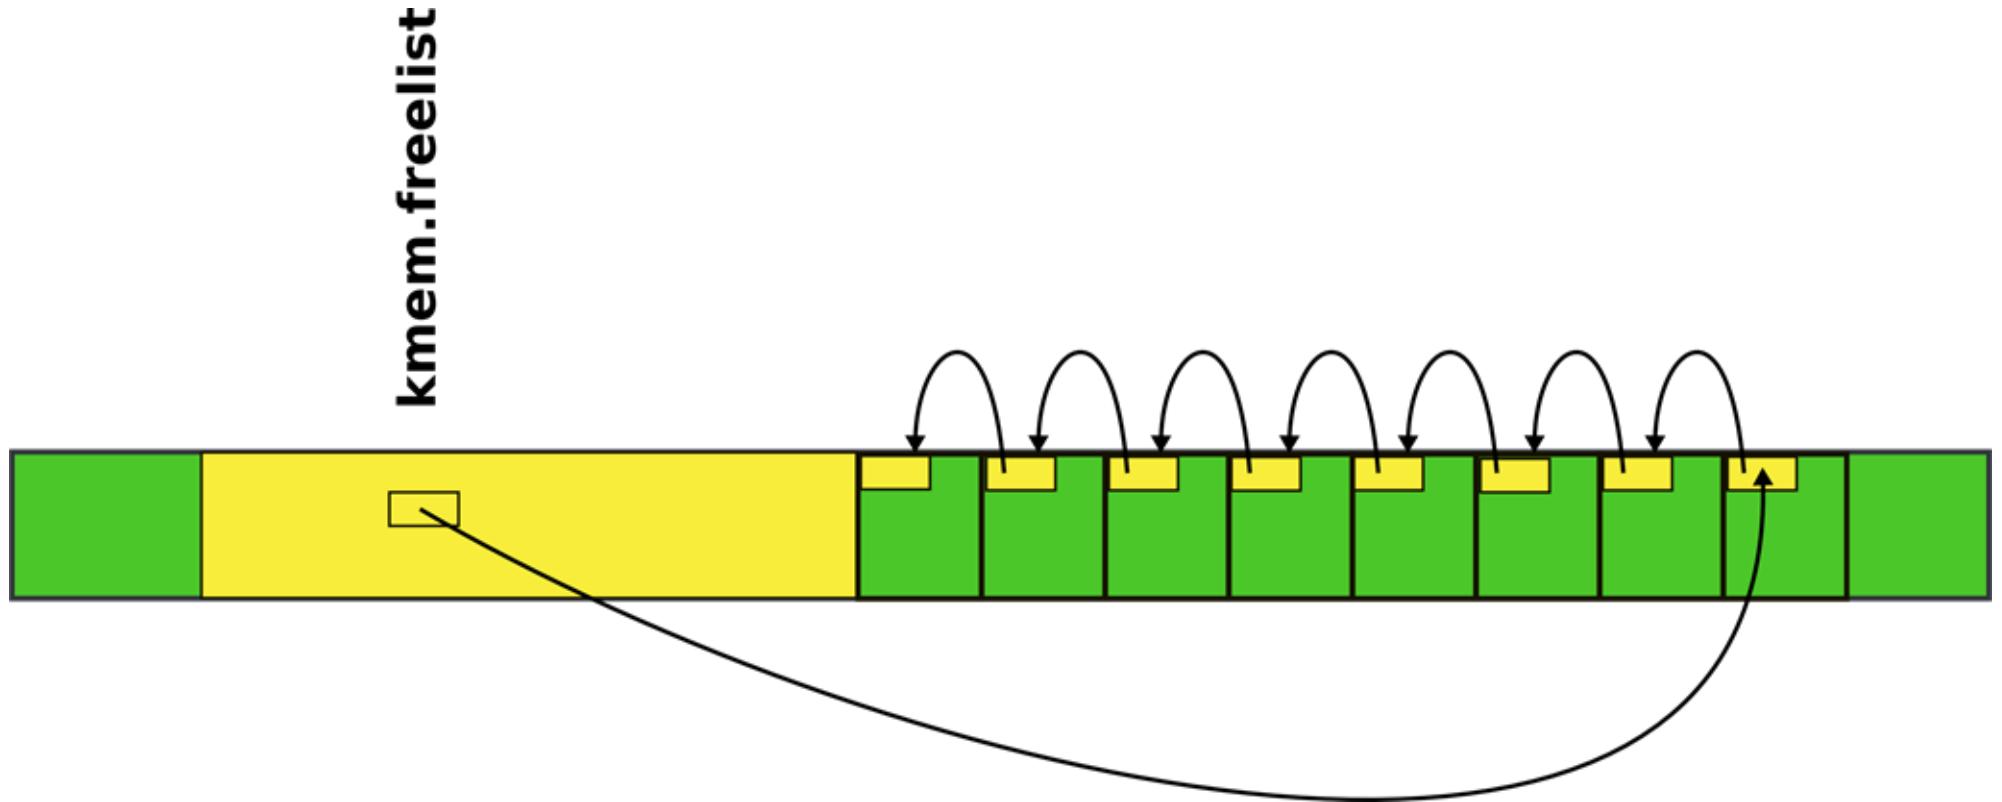
```
3077 r = (struct run*)v;
```

```
3078 r->next = kmem.freelist;
```

```
3079 kmem.freelist = r;
```

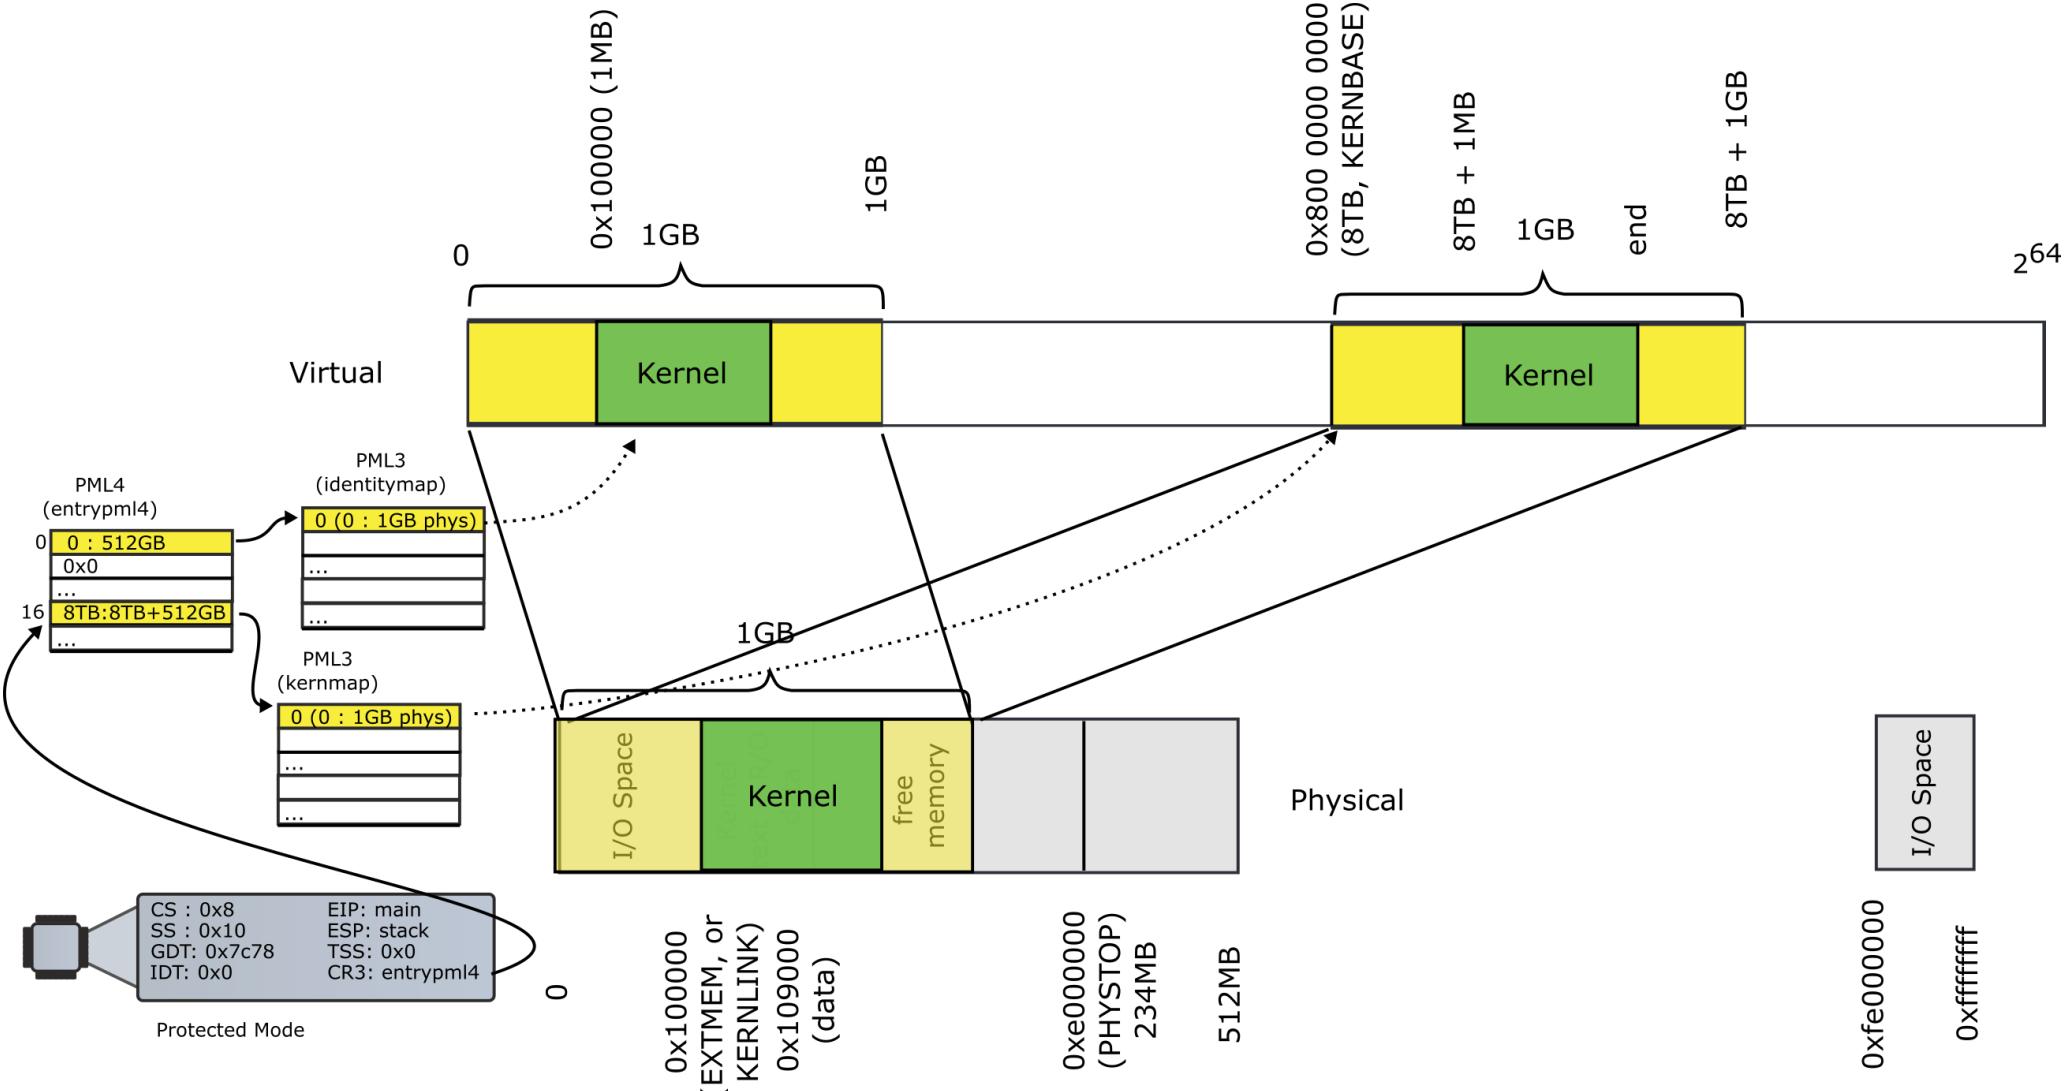
```
...
```

```
2832 }
```

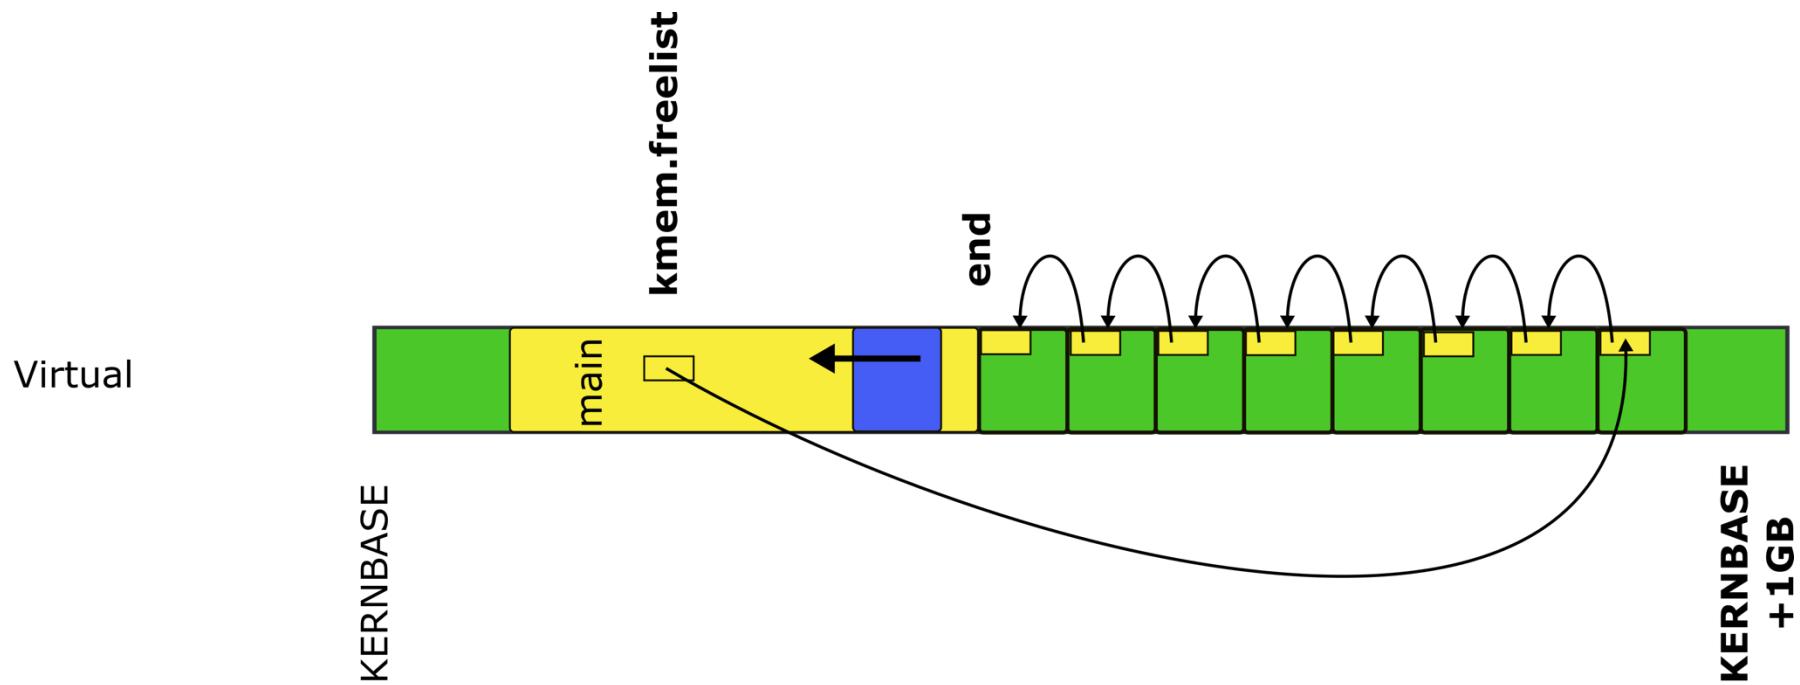


- Where can we get memory to keep the list itself?
- Note, the list is maintained within each page
- It has to write each page though to update the “next” pointer

# There is free memory in the 1GB page we've mapped



# Donate this free memory to the allocator



- Take memory from the end of the kernel binary
- To the end of the 4MB page

# kinit1(): initialize the allocator with free memory

```
1366 int
1367 main(void)
1368 {
1369     void *kinit1_end = (1024*1024*1024 > PHYSTOP) ? P2V(PHYSTOP)
1370                 : P2V(1024*1024*1024);
1371     kinit1(end, kinit1_end); // phys page allocator
1372     kvmalloc(); // kernel page table
1373     mpinit(); // detect other processors
1374     lapicinit(); // interrupt controller
1375     seginit(); // segment descriptors
1376     picinit(); // disable pic
1377     ioapicinit(); // another interrupt controller
```

...

# Freerange()

```
3030 kinit1(void *vstart, void *vend)
```

```
3031 {
```

```
...
```

```
3034 freerange(vstart, vend);
```

```
3035 }
```

- Free range of memory from `vstart` to `vend` giving it to the allocator
- i.e., adding pages to the list

# freerange()

```
3051 freerange(void *vstart, void *vend)
3052 {
3053     char *p;
3054     p = (char*)PGROUNDUP((uint)vstart);
3055     for(; p + PGSIZE <= (char*)vend; p += PGSIZE)
3056         kfree(p);
3057 }
```

- freerange() internally simply frees the pages from vstart to vend
- kfree() adds them to the allocator list

# Where do we start?

```
1366 int  
1367 main(void)  
1368 {  
1369     void *kinit1_end = (1024*1024*1024 > PHYSTOP) ? P2V(PHYSTOP)  
1370                 : P2V(1024*1024*1024);  
1371     kinit1(end, kinit1_end); // phys page allocator
```

- What is this **end**?

```
1311 extern char end[];
```

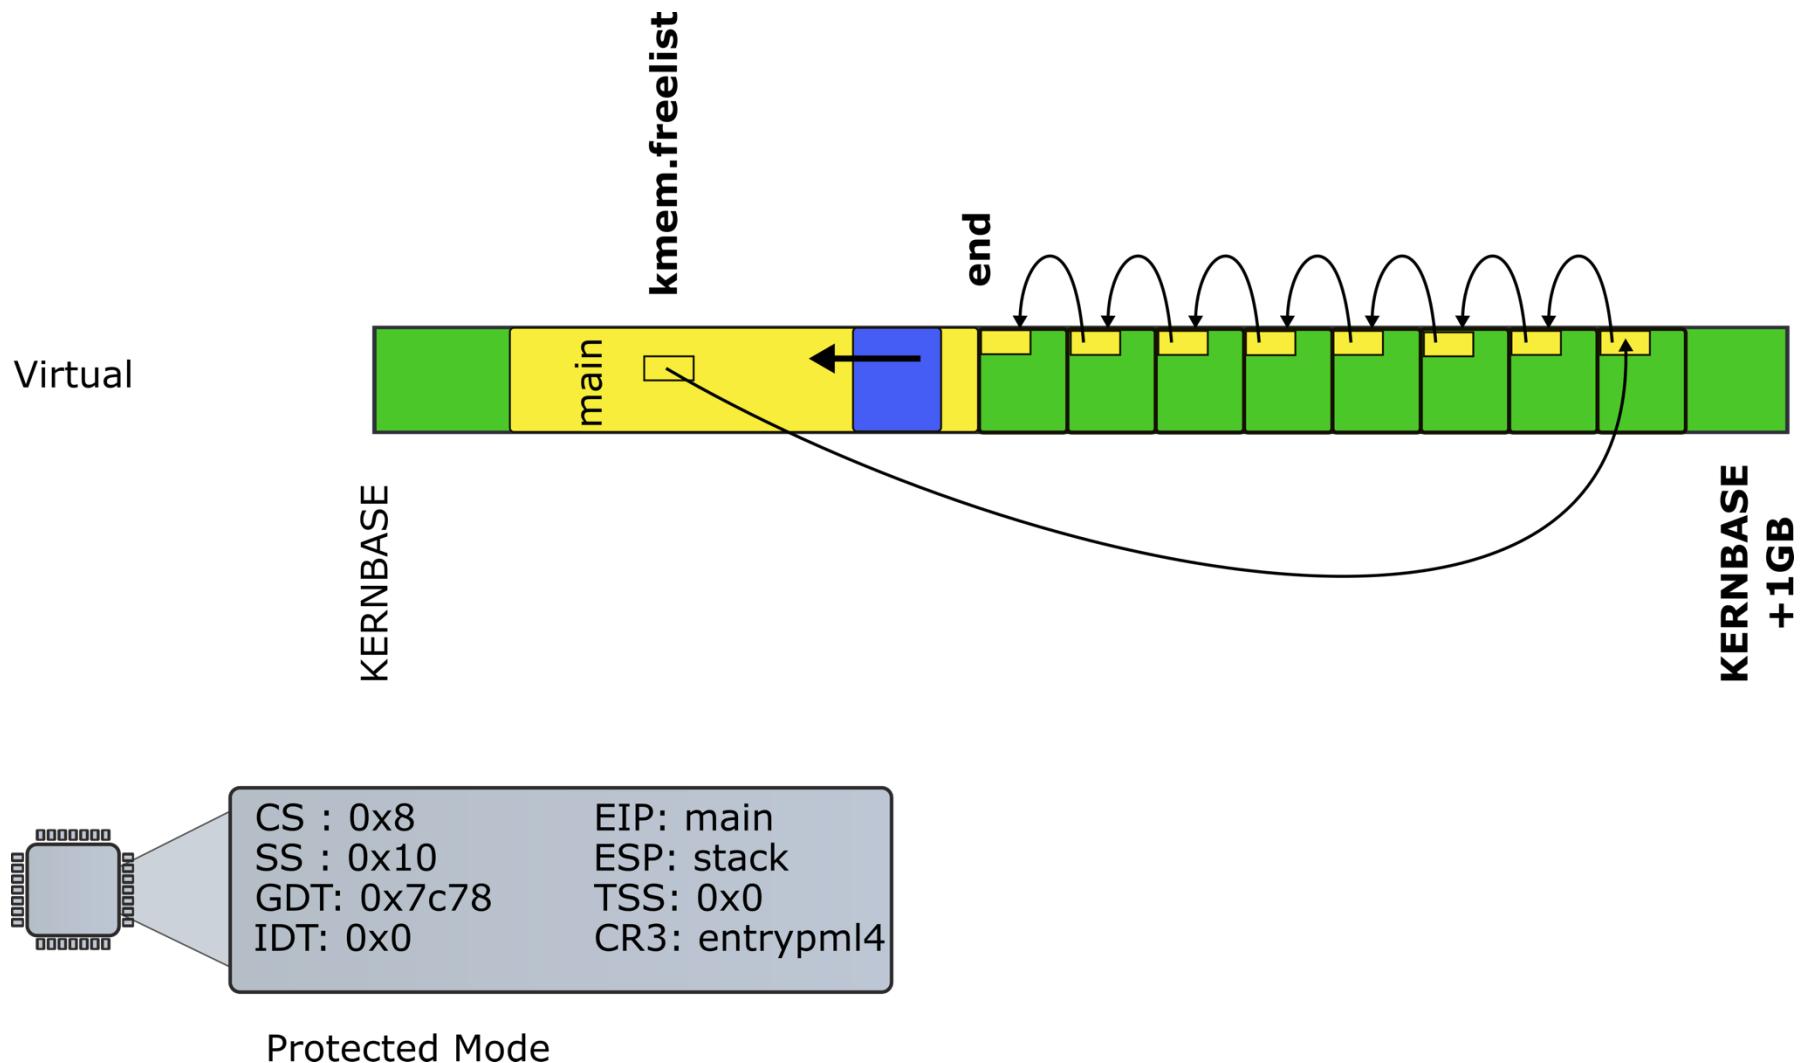
# Where do we start?

```
1366 int  
1367 main(void)  
1368 {  
1369     void *kinit1_end = (1024*1024*1024 > PHYSTOP) ? P2V(PHYSTOP)  
1370                 : P2V(1024*1024*1024);  
1371     kinit1(end, kinit1_end); // phys page allocator
```

- What is this **end**?

```
1311 extern char end[]; // first address after  
                           kernel loaded from ELF file
```

# Donate this free memory to the allocator



# Recap

- Kernel has a memory allocator
- It allocates memory in chunks of 4KB
- Good enough to maintain kernel data structures


# Kernel page table (for 4KB page tables)

# Back to main(): Kernel address space

```
1366 int
1367 main(void)
1368 {
1369     void *kinit1_end = (1024*1024*1024 > PHYSTOP) ? P2V(PHYSTOP)
1370                 : P2V(1024*1024*1024);
1371     kinit1(end, kinit1_end); // phys page allocator
1372     kvmalloc(); // kernel page table
1373     mpinit(); // detect other processors
1374     lapicinit(); // interrupt controller
1375     seginit(); // segment descriptors
1376     picinit(); // disable pic
1377     ioapicinit(); // another interrupt controller
```

- What do you think has to happen?
  - i.e., how do we construct a kernel page table?


# Recap: our goal




# Outline

- Map a region of virtual memory into page tables
  - Start from 2GBs
  - Iterate memory page by page
  - Allocate page table directory and page tables as we go
  - Fill in page table entries with proper physical addresses


# Allocate page table directory entry




# Allocate next level page table




# Locate PTE entry



# Update mapping with physical addr



# Move to next page



This is exactly what kernel is doing  
(let's read the [source code](#))

```
1316 int
1317 main(void)
1318 {
1319     kinit1(end, P2V(4*1024*1024)); // phys page allocator
1320     kvmalloc(); // kernel page table
1321     mpinit(); // detect other processors
1322     lapicinit(); // interrupt controller
1323     seginit(); // segment descriptors
1324     cprintf("\ncpu%d: starting xv6\n\n", cpunum());
1325     picinit(); // another interrupt controller
1326     ioapicinit(); // another interrupt controller
1327     consoleinit(); // console hardware
1328     uartinit(); // serial port
...
1340 }
```

# Allocate page tables

# kvmalloc()


```
1857 kvmalloc(void)
```

```
1858 {
```

```
1859   kpgdir = setupkvm();
```


```
1860   switchkvm();
```

```
1861 }
```




```
1836 pde_t*
1837 setupkvm(void)
1838 {
1839     pde_t *pgdir;
1840     struct kmap *k;
1841
1842     if((pgdir = (pde_t*)kalloc()) == 0)
1843         return 0;
1844     memset(pgdir, 0, PGSIZE);
...
1847     for(k = kmap; k < &kmap[NELEM(kmap)]; k++)
1848         if(mappages(pgdir, k->virt, k->phys_end - k->phys_start,
1849                     (uint)k->phys_start, k->perm) < 0)
1850             return 0;
1851     return pgdir;
1852 }
```

# Allocate page table directory




```
1836 pde_t*
1837 setupkvm(void)
1838 {
1839     pde_t *pgdir;
1840     struct kmap *k;
1841
1842     if((pgdir = (pde_t*)kalloc()) == 0)
1843         return 0;
1844     memset(pgdir, 0, PGSIZE);
...
1847     for(k = kmap; k < &kmap[NELEM(kmap)]; k++)
1848         if(mappages(pgdir, k->virt, k->phys_end - k->phys_start,
1849                     (uint)k->phys_start, k->perm) < 0)
1850             return 0;
1851     return pgdir;
1852 }
```

What is the address of this page?




```
1836 pde_t*
1837 setupkvm(void)
1838 {
1839     pde_t *pgdir;
1840     struct kmap *k;
1841
1842     if((pgdir = (pde_t*)kalloc()) == 0)
1843         return 0;
1844     memset(pgdir,
...
1847     for(k = kmap; k < mappages;
1848         if(mappages(
1849             (uint)k-
1850         return 0;
1851     return pgdir;
1852 }
```

What is the address of this page?




```
1836 pde_t*
1837 setupkvm(void)
1838 {
1839     pde_t *pgdir;
1840     struct kmap *k;
1841
1842     if((pgdir = (pde_t*)kalloc()) == 0)
1843         return 0;
1844     memset(pgdir, 0, PGSIZE);
...
1847     for(k = kmap; k < &kmap[NELEM(kmap)]; k++)
1848         if(mappages(pgdir, k->virt, k->phys_end - k->phys_start,
1849                     (uint)k->phys_start, k->perm) < 0)
1850             return 0;
1851     return pgdir;
1852 }
```

What is the address of this page?



# Allocate page table directory




```

1836 pde_t*
1887 setupkvm(void)
1838 {
1839     pde_t *pgdir;
1840     struct kmap *k;
1841
1842     if((pgdir = (pde_t*)kalloc()) == 0)
1843         return 0;
1844     memset(pgdir, 0, PGSIZE);
...
1847     for(k = kmap; k < &kmap[NELEM(kmap)]; k++)
1848         if(mappages(pgdir, k->virt, k->phys_end - k->phys_start,
1849                     (uint)k->phys_start, k->perm) < 0)
1850             return 0;
1851     return pgdir;
1852 }

```

# Iterate in a loop: map physical pages




```


1836 pde_t*
1887 setupkvm(void)
1838 {
1839     pde_t *pgdir;
1840     struct kmap *k;
1841
1842     if((pgdir = (pde_t*)kalloc()) == 0)
1843         return 0;
1844     memset(pgdir, 0, PGSIZE);
...
1847     for(k = kmap; k < &kmap[NELEM(kmap)]; k++)
1848         if(mappages(pgdir, k->virt, k->phys_end - k->phys_start,
1849                     (uint)k->phys_start, k->perm) < 0)
1850             return 0;
1851     return pgdir;
1852 }

```

# Iterate in a loop: map physical pages




# Kernel map



# Kernel map

[PollEv.com/antonburtsev](http://PollEv.com/antonburtsev)



# Kmap – kernel map

```
1823 static struct kmap {
```

```
1824     void *virt;
```

Physical

```
1825     uint phys_start;
```



```
1826     uint phys_end;
```

```
1827     int perm;
```

```
1828 } kmap[] = {
```

```
1829     { (void*)KERNBASE, 0, EXTMEM, PTE_W}, // I/O space
```

```
1830     { (void*)KERNLINK, V2P(KERNLINK), V2P(data), 0}, //text+rodata
```

```
1831     { (void*)data, V2P(data), PHYSTOP, PTE_W}, // kern data+memory
```

```
1832     { (void*)DEVSSPACE, DEVSSPACE, 0, PTE_W}, // more devices
```


```
1833 };
```

```

1836 pde_t*
1887 setupkvm(void)
1838 {
1839     pde_t *pgdir;
1840     struct kmap *k;
1841
1842     if((pgdir = (pde_t*)kalloc()) == 0)
1843         return 0;
1844     memset(pgdir, 0, PGSIZE);
...
1847     for(k = kmap; k < &kmap[NELEM(kmap)]; k++)
1848         if(mappages(pgdir, k->virt, k->phys_end - k->phys_start,
1849                     (uint)k->phys_start, k->perm) < 0)
1850             return 0;
1851     return pgdir;
1852 }

```

# Start virtual address



```
1836 pde_t*
```

```
1887 setupkvm(void)
```

```
1838 {
```

```
1839 pde_t *pgdir;
```

```
1840 struct kmap *k;
```

```
1841
```

```
1842 if((pgdir = (pde_t*)kalloc()) == 0)
```

```
1843     return 0;
```

```
1844 memset(pgdir, 0, PGSIZE);
```

```
...
```

```
1847 for(k = kmap; k < &kmap[NELEM(kmap)]; k++)
```

```
1848     if(mappages(pgdir, k->virt, k->phys_end - k->phys_start,
```


```
1849             (uint)k->phys_start, k->perm) < 0)
```

```
1850     return 0;
```

```
1851     return pgdir;
```

```
1852 }
```

# Size




```

1836 pde_t*
1887 setupkvm(void)
1838 {
1839     pde_t *pgdir;
1840     struct kmap *k;
1841
1842     if((pgdir = (pde_t*)kalloc()) == 0)
1843         return 0;
1844     memset(pgdir, 0, PGSIZE);
...
1847     for(k = kmap; k < &kmap[NELEM(kmap)]; k++)
1848         if(mappages(pgdir, k->virt, k->phys_end - k->phys_start,
1849                     (uint)k->phys_start, k->perm) < 0)
1850             return 0;
1851     return pgdir;
1852 }

```


# Start physical address



```

1779 mappages(pde_t *pgdir, void *va, uint size, uint pa, int perm)
1780 {
1781     char *a, *last;
1782     pte_t *pte;
1783
1784     a = (char*)PGROUNDDOWN((uint)va);
1785     last = (char*)PGROUNDDOWN(((uint)va) + size - 1);
1786     for(;;){
1787         if((pte = walkpgdir(pgdir, a, 1)) == 0)
1788             return -1;
1789         if(*pte & PTE_P)
1790             panic("remap");
1791         *pte = pa | perm | PTE_P;
1792         if(a == last)
1793             break;
1794         a += PGSIZE;
1795         pa += PGSIZE;
1796     }
1797     return 0;
1798 }
```


# Inside mappages()



- Get the start (`a`) and end (`last`) pages fo the virtual address range we are mapping
- Then work in a loop mapping every page one by one


```

1779 mappages(pde_t *pgdir, void *va, uint size, uint pa, int perm)
1780 {
1781     char *a, *last;
1782     pte_t *pte;
1783
1784     a = (char*)PGROUNDDOWN((uint)va);
1785     last = (char*)PGROUNDDOWN(((uint)va) + size - 1);
1786     for(;;){
1787         if((pte = walkpgdir(pgdir, a, 1)) == 0)
1788             return -1;
1789         if(*pte & PTE_P)
1790             panic("remap");
1791         *pte = pa | perm | PTE_P;
1792         if(a == last)
1793             break;
1794         a += PGSIZE;
1795         pa += PGSIZE;
1796     }
1797     return 0;
1798 }
```




- First lookup the page table entry (pte) corresponding to the virtual address (a) we're mapping

# Locate the page table entry




```

1779 mappages(pde_t *pgdir, void *va, uint size, uint pa, int perm)
1780 {
1781     char *a, *last;
1782     pte_t *pte;
1783
1784     a = (char*)PGROUNDDOWN((uint)va);
1785     last = (char*)PGROUNDDOWN(((uint)va) + size - 1);
1786     for(;;){
1787         if((pte = walkpgdir(pgdir, a, 1)) == 0)
1788             return -1;
1789         if(*pte & PTE_P)
1790             panic("remap");
1791         *pte = pa | perm | PTE_P;
1792         if(a == last)
1793             break;
1794         a += PGSIZE;
1795         pa += PGSIZE;
1796     }
1797     return 0;
1798 }
```




- Update the page directory entry (\*pte) with the physical address (pa)

# Update mapping with physical addr



```

1779 mappages(pde_t *pgdir, void *va, uint size, uint pa, int perm)
1780 {
1781     char *a, *last;
1782     pte_t *pte;
1783
1784     a = (char*)PGROUNDDOWN((uint)va);
1785     last = (char*)PGROUNDDOWN(((uint)va) + size - 1);
1786     for(;;){
1787         if((pte = walkpgdir(pgdir, a, 1)) == 0)
1788             return -1;
1789         if(*pte & PTE_P)
1790             panic("remap");
1791         *pte = pa | perm | PTE_P;
1792         if(a == last)
1793             break;
1794         a += PGSIZE;
1795         pa += PGSIZE;
1796     }
1797     return 0;
1798 }
```




- But we need a function that locates the pte for us...


# What should it look like?

- A function takes a virtual address
- Returns a page table directory entry that maps it

# Recap of the page table



# Locate the PDE frist



```
1754 walkpgdir(pde_t *pgdir, const void *va, int alloc)
1755 {
1756     pde_t *pde;
1757     pte_t *pgtab;
1758
1759     pde = &pgdir[PDX(va)];
1760     if(*pde & PTE_P){
1761         pgtab = (pte_t*)P2V(PTE_ADDR(*pde));
1762     } else {
1763         if(!alloc || (pgtab = (pte_t*)kalloc()) == 0)
1764             return 0;
1765         // Make sure all those PTE_P bits are zero.
1766         memset(pgtab, 0, PGSIZE);
1767
1768         ...
1769         *pde = V2P(pgtab) | PTE_P | PTE_W | PTE_U;
1770     }
1771     return &pgtab[PTX(va)];
1772 }
1773 }
```

## walkpgdir(): walk page table

- Locate the page directory entry (\*pde)

# PDX()

```
0855 // +-----10-----+-----10-----+-----12-----+
0856 // | Page Directory | Page Table | Offset within Page |
0857 // | Index       | Index       |           |
0858 // +-----+-----+-----+
0859 // \--- PDX(va) ---/ \--- PTX(va) ---/
0860
```

```
0861 // page directory index
0862 #define PDX(va) (((uint)(va) >> PDXSHIFT) & 0x3FF)
0863
0864 // page table index
0865 #define PTX(va) (((uint)(va) >> PTXSHIFT) & 0x3FF)
```

...

```
0876 #define PTXSHIFT 12 // offset of PTX in a linear address
0877 #define PDXSHIFT 22 // offset of PDX in a linear address
```

# PDX()

```
0855 // +-----10-----+-----10-----+-----12-----+
```

```
0856 // | Page Directory | Page Table | Offset within Page |
```

```
0857 // | Index | Index | |
```

```
0858 // +-----+-----+-----+-----+
```

```
0859 // \--- PDX(va) --/ \--- PTX(va) --/
```

```
0860
```

```
0861 // page directory index
```

```
0862 #define PDX(va) (((uint)(va) >> PDXSHIFT) & 0x3FF)
```

```
0863
```

```
0864 // page table index
```

```
0865 #define PTX(va) (((uint)(va) >> PTXSHIFT) & 0x3FF)
```

```
...
```

```
0876 #define PTXSHIFT 12 // offset of PTX in a linear address
```


```
0877 #define PDXSHIFT 22 // offset of PDX in a linear address
```

```
1754 walkpgdir(pde_t *pgdir, const void *va, int alloc)
1755 {
1756     pde_t *pde;
1757     pte_t *pgtab;
1758
1759     pde = &pgdir[PDX(va)];
1760     if(*pde & PTE_P){
1761         pgtab = (pte_t*)P2V(PTE_ADDR(*pde));
1762     } else {
1763         if(!alloc || (pgtab = (pte_t*)kalloc()) == 0)
1764             return 0;
1765         // Make sure all those PTE_P bits are zero.
1766         memset(pgtab, 0, PGSIZE);
1767
1768         ...
1769         *pde = V2P(pgtab) | PTE_P | PTE_W | PTE_U;
1770     }
1771     return &pgtab[PTX(va)];
1772 }
1773 }
```

## walkpgdir(): walk page table

- Check if page table is allocated (present)

# Check if level 2 page table is allocated




```
1754 walkpgdir(pde_t *pgdir, const void *va, int alloc)
1755 {
1756     pde_t *pde;
1757     pte_t *pgtab;
1758
1759     pde = &pgdir[PDX(va)];
1760     if(*pde & PTE_P){
1761         pgtab = (pte_t*)P2V(PTE_ADDR(*pde));
1762     } else {
1763         if(!alloc || (pgtab = (pte_t*)kalloc()) == 0)
1764             return 0;
1765         // Make sure all those PTE_P bits are zero.
1766         memset(pgtab, 0, PGSIZE);
1767         ...
1768         *pde = V2P(pgtab) | PTE_P | PTE_W | PTE_U;
1769     }
1770     return &pgtab[PTX(va)];
1771 }
1772 }
```

walkpgdir(): walk page table

- Allocate if needed

# See if the next page table level exists (PTE\_P is set)




```
1754 walkpgdir(pde_t *pgdir, const void *va, int alloc)
1755 {
1756     pde_t *pde;
1757     pte_t *pgtab;
1758
1759     pde = &pgdir[PDX(va)];
1760     if(*pde & PTE_P){
1761         pgtab = (pte_t*)P2V(PTE_ADDR(*pde));
1762     } else {
1763         if(!alloc || (pgtab = (pte_t*)kalloc()) == 0)
1764             return 0;
1765         // Make sure all those PTE_P bits are zero.
1766         memset(pgtab, 0, PGSIZE);
1767
1768         ...
1769         *pde = V2P(pgtab) | PTE_P | PTE_W | PTE_U;
1770     }
1771     return &pgtab[PTX(va)];
1772 }
1773 }
```

## walkpgdir(): walk page table

- If exists, get the address of the next level

# PDE contains 20 bits which represent physical page number



# Getting level 2 page

```
1761 pgtab = (pte_t*)P2V(PTE_ADDR(*pde));
```

- We need two things
- Convert from 20 bits of physical page number to physical address of the page
- PTE\_ADDR(\*pde)
- Convert from physical address of that page to virtual address
- P2V(...)
  - We can't access physical addresses directly
  - We can only access virtual addresses
  - Registers, mov instructions, etc. contain virtual addresses
  - Physical address have to be mapped by the current page table


# Step 1

- Convert from 20 bits of physical page number to physical address of the page
- `PTE_ADDR(*pde)`
- This is trivial

# Step 2

- Convert from physical address of that page to virtual address
- P2V(...)
- This seems a bit tricky


# Remember how we mapped the kernel?



```
0207 #define KERNBASE 0x80000000 // First kernel virtual address
```

```
0210 #define V2P(a) (((uint) (a)) - KERNBASE)
```

```
0211 #define P2V(a) (((void *) (a)) + KERNBASE)
```




```
1754 walkpgdir(pde_t *pgdir, const void *va, int alloc)
1755 {
1756     pde_t *pde;
1757     pte_t *pgtab;
1758
1759     pde = &pgdir[PDX(va)];
1760     if(*pde & PTE_P){
1761         pgtab = (pte_t*)P2V(PTE_ADDR(*pde));
1762     } else {
1763         if(!alloc || (pgtab = (pte_t*)kalloc()) == 0)
1764             return 0;
1765         // Make sure all those PTE_P bits are zero.
1766         memset(pgtab, 0, PGSIZE);
1767
1768         ...
1769         *pde = V2P(pgtab) | PTE_P | PTE_W | PTE_U;
1770     }
1771     return &pgtab[PTX(va)];
1772 }
1773 }
```


## walkpgdir(): walk page table

- Return pointer to the PTE

# Poll: what are we returning?




# Return a pointer to PTE



Back to mappages() function that maps a region of virtual memory into continuous region of physical memory

```


1779 mappages(pde_t *pgdir, void *va, uint size, uint pa, int perm)
1780 {
1781     char *a, *last;
1782     pte_t *pte;
1783
1784     a = (char*)PGROUNDDOWN((uint)va);
1785     last = (char*)PGROUNDDOWN(((uint)va) + size - 1);
1786     for(;;){
1787         if((pte = walkpgdir(pgdir, a, 1)) == 0)
1788             return -1;
1789         if(*pte & PTE_P)
1790             panic("remap");
1791         *pte = pa | perm | PTE_P;
1792         if(a == last)
1793             break;
1794         a += PGSIZE;
1795         pa += PGSIZE;
1796     }
1797     return 0;
1798 }
```



Remember we just  
discussed [walkpgdir\(\)](#)

```


1779 mappages(pde_t *pgdir, void *va, uint size, uint pa, int perm)
1780 {
1781     char *a, *last;
1782     pte_t *pte;
1783
1784     a = (char*)PGROUNDDOWN((uint)va);
1785     last = (char*)PGROUNDDOWN(((uint)va) + size - 1);
1786     for(;;){
1787         if((pte = walkpgdir(pgdir, a, 1)) == 0)
1788             return -1;
1789         if(*pte & PTE_P)
1790             panic("remap");
1791         *pte = pa | perm | PTE_P;
1792         if(a == last)
1793             break;
1794         a += PGSIZE;
1795         pa += PGSIZE;
1796     }
1797     return 0;
1798 }
```



# Page present (PTE\_P) – panic

```


1779 mappages(pde_t *pgdir, void *va, uint size, uint pa, int perm)
1780 {
1781     char *a, *last;
1782     pte_t *pte;
1783
1784     a = (char*)PGROUNDDOWN((uint)va);
1785     last = (char*)PGROUNDDOWN(((uint)va) + size - 1);
1786     for(;;){
1787         if((pte = walkpgdir(pgdir, a, 1)) == 0)
1788             return -1;
1789         if(*pte & PTE_P)
1790             panic("remap");
1791         *pte = pa | perm | PTE_P;
1792         if(a == last)
1793             break;
1794         a += PGSIZE;
1795         pa += PGSIZE;
1796     }
1797     return 0;
1798 }
```



- Update page table entry
- Where does **\*pte** point?
  - pa – physical address of the page

```

1779 mappages(pde_t *pgdir, void *va, uint size, uint pa, int perm)
1780 {
1781     char *a, *last;
1782     pte_t *pte;
1783
1784     a = (char*)PGROUNDDOWN((uint)va);
1785     last = (char*)PGROUNDDOWN(((uint)va) + size - 1);
1786     for(;;){
1787         if((pte = walkpgdir(pgdir, a, 1)) == 0)
1788             return -1;
1789         if(*pte & PTE_P)
1790             panic("remap");
1791         *pte = pa | perm | PTE_P;
1792         if(a == last)
1793             break;
1794         a += PGSIZE;
1795         pa += PGSIZE;
1796     }
1797     return 0;
1798 }
```



- Move to the next page

# kvmalloc()


```
1757 kvmalloc(void)
```

```
1758 {
```

```
1759   kpgdir = setupkvm();
```

```
1760   switchkvm();
```

```
1761 }
```



# Switch to the new page table


```
1765 void
```

```
1766 switchkvm(void)
```

```
1767 {
```

```
1768 lcr3(v2p(kpgdir));
```

```
1769 }
```



# Recap

- Kernel has a memory allocator
- Kernel has its own address space
- It uses 4KB page tables
- It is ready to create processes

```
1317 main(void)
1318 {
1319     kinit1(end, P2V(4*1024*1024)); // phys page allocator
1320     kvmalloc(); // kernel page table
1321     mpinit(); // detect other processors
1322     lapicinit(); // interrupt controller
1323     seginit(); // segment descriptors
1324     cprintf("\ncpu%d: starting xv6\n\n", cpunum());
1325     picinit(); // another interrupt controller
1326     ioapicinit(); // another interrupt controller
1327     consoleinit(); // console hardware
1328     uartinit(); // serial port
1329     pinit(); // process table
1330     tvinit(); // trap vectors
1331     binit(); // buffer cache
1332     fileinit(); // file table
1333     ideinit(); // disk
1334     if(!ismp)
1335     timerinit(); // uniprocessor timer
1336     startothers(); // start other processors
1337     kinit2(P2V(4*1024*1024), P2V(PHYSTOP)); // must come after startothers()
```

main()

```
1317 main(void)
1318 {
1319     kinit1(end, P2V(4*1024*1024)); // phys page allocator
1320     kvmalloc(); // kernel page table
1321     mpinit(); // detect other processors
1322     lapicinit(); // interrupt controller
1323     seginit(); // segment descriptors
1324     cprintf("\ncpu%d: starting xv6\n\n", cpunum());
1325     picinit(); // another interrupt controller
1326     ioapicinit(); // another interrupt controller
1327     consoleinit(); // console hardware
1328     uartinit(); // serial port
1329     pinit(); // process table
1330     tvinit(); // trap vectors
1331     binit(); // buffer cache
1332     fileinit(); // file table
1333     ideinit(); // disk
1334     if(!ismp)
1335         timerinit(); // uniprocessor timer
1336     startothers(); // start other processors
1337     kinit2(P2V(4*1024*1024), P2V(PHYSTOP)); // must come after startothers()
1338     userinit(); // first user process
1339     mpmain(); // finish this processor's setup
1340 }
```

main()

# Initialize GDT

```
1712 // Set up CPU's kernel segment descriptors.
```

```
1713 // Run once on entry on each CPU.
```

```
1714 void
```

```
1715 seginit(void)
```

```
1716 {
```

```
1717 struct cpu *c;
```

```
1718
```

```
1719 // Map "logical" addresses to virtual addresses using identity map.
```

```
1720 // Cannot share a CODE descriptor for both kernel and user
```

```
1721 // because it would have to have DPL_USR, but the CPU forbids
```

```
1722 // an interrupt from CPL=0 to DPL=3.
```

```
1723 c = &cpus[cpuid()];
```

```
1724 c->gdt[SEG_KCODE] = SEG(STA_X|STA_R, 0, 0xffffffff, 0);
```


```
1725 c->gdt[SEG_KDATA] = SEG(STA_W, 0, 0xffffffff, 0);
```

```
1726 c->gdt[SEG_UCODE] = SEG(STA_X|STA_R, 0, 0xffffffff, DPL_USER);
```

```
1727 c->gdt[SEG_UDATA] = SEG(STA_W, 0, 0xffffffff, DPL_USER);
```

```
1728 lgdt(c->gdt, sizeof(c->gdt));
```


```
1729 }
```



# Struct CPU

```
2300 // Per-CPU state  
2301 struct cpu {  
2302     uchar apicid;          // Local APIC ID  
2303     struct context *scheduler; // swtch() here to enter scheduler  
2304     struct taskstate ts; // Used by x86 to find stack for interrupt  
2305     struct segdesc gdt[NSEGGS]; // x86 global descriptor table  
2306     volatile uint started;    // Has the CPU started?  
2307     int ncli;                // Depth of pushcli nesting.  
2308     int intena;              // Were interrupts enabled before pushcli?  
2309     struct proc *proc;        // The process running on this cpu or null  
2310 };  
2311  
2312 extern struct cpu cpus[NCPU];
```

# Segment descriptor



L — 64-bit code segment (IA-32e mode only)

AVL — Available for use by system software

BASE — Segment base address

D/B — Default operation size (0 = 16-bit segment; 1 = 32-bit segment)

DPL — Descriptor privilege level

G — Granularity

LIMIT — Segment Limit

P — Segment present

S — Descriptor type (0 = system; 1 = code or data)

TYPE — Segment type

# Segment Descriptor

```
0724 // Segment Descriptor  
0725 struct segdesc {  
0726     uint lim_15_0 : 16; // Low bits of segment limit  
0727     uint base_15_0 : 16; // Low bits of segment base address  
0728     uint base_23_16 : 8; // Middle bits of segment base address  
0729     uint type : 4;    // Segment type (see STS_constants)  
0730     uint s : 1;      // 0 = system, 1 = application  
0731     uint dpl : 2;    // Descriptor Privilege Level  
0732     uint p : 1;      // Present  
0733     uint lim_19_16 : 4; // High bits of segment limit  
0734     uint avl : 1;    // Unused (available for software use)  
0735     uint rsv1 : 1;   // Reserved  
0736     uint db : 1;    // 0 = 16-bit segment, 1 = 32-bit segment  
0737     uint g : 1;     // Granularity: limit scaled by 4K when set  
0738     uint base_31_24 : 8; // High bits of segment base address  
0739 };
```

Thank you!

(Next time: interrupts!)